

Interpretable Strategy Synthesis for Competitive Games

Thesis Defense Presentation

Abhijeet Krishnan

Department of Computer Science
North Carolina State University

July 16, 2024

Publications

- Krishnan, Abhijeet, Colin M. Potts, Arnav Jhala, Harshad Khadilkar, Shirish Karande and Chris Martens. "Learning Explainable Representations of Complex Game-playing Strategies." *Proceedings of the Eleventh Annual Conference on Advances in Cognitive Systems*. 2024.
- Villalobos-Arias, Leonardo, Derek Martin, Abhijeet Krishnan, Madeleine Gagné, Colin M. Potts and Arnav Jhala. "Modeling Risk in Reinforcement Learning: A Literature Mapping." *arXiv preprint arXiv:2312.05231*. 2023.
- Krishnan, Abhijeet and Chris Martens. "Synthesizing Chess Tactics from Player Games." In *Workshop on Artificial Intelligence for Strategy Games (SG) and Esports Analytics (EA), 18th AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*. 2022.
- Krishnan, Abhijeet and Chris Martens. "Towards the Automatic Synthesis of Interpretable Chess Tactics." In *Explainable Agency in Artificial Intelligence Workshop, 36th AAAI Conference on Artificial Intelligence*. 2022.
- Krishnan, Abhijeet, Aaron Williams, and Chris Martens. "Towards Action Model Learning for Player Modeling." *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*. Vol. 16. No. 1. 2020.
- Krishnan, Abhijeet and Chris Martens. "Rule-based Cognitive Modeling via Human-Computer Interaction." Poster presented at: *5th LAS Research Symposium*; 2019 Dec 10; Raleigh, NC.

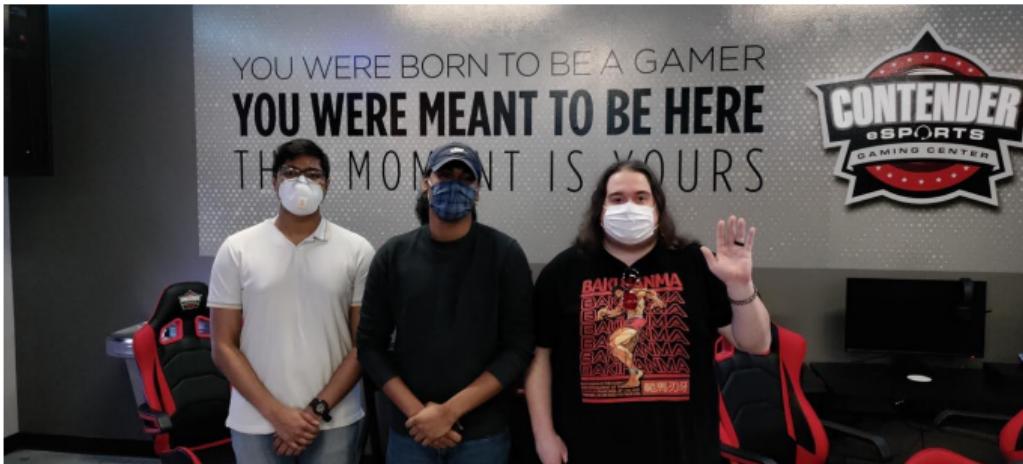
Getting Into Competitive Gaming



Kevin Brown

Participating in a competitive *Tekken 7* tournament

Wanting to Improve



Contender Esports

Placing 3rd

20 FEB 2022
Tekken 7 at Tournament of Thrones 18
Raleigh, NC 11 Entrants

2nd
PLACE
out of 11

1. HTL Clam
2. Abhijeet
3. Ingress

15 JAN 2022
Tekken 7 at Last Level: Raleigh 5
Cary, NC 8 Entrants

3rd
PLACE
out of 8

1. HTL Clam
2. DeL
3. Abhijeet

21 NOV 2021
Tekken 7 at Tournament of Thrones 15
Raleigh, NC 6 Entrants

2nd
PLACE
out of 6

1. HTL Clam
2. STFU Abhijeet
3. Ahematical

start.gg

Results of recent tournaments

How to Improve?

Improving at a skill —

How to Improve?

Improving at a skill —

- Play more

How to Improve?

Improving at a skill —

- Play more → *time-consuming*

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach → *not structured*

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach → *not structured*
- Talk to players

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach → *not structured*
- Talk to players → *not scalable*

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach → *not structured*
- Talk to players → *not scalable*
- Study better players

How to Improve?

Improving at a skill —

- Play more → *time-consuming*
- Use a learning resource → *not standardized*
- Find a coach → *not structured*
- Talk to players → *not scalable*
- Study better players → *difficult*

Possible Solution

Could we *automatically* generate *explainable* strategies for players to learn from?

Strategies are Mental Models of Games

Claim

Strategies can be treated as mental models of games

Strategies are Mental Models of Games

Claim

Strategies can be treated as mental models of games

- Players develop *mental models* of games to explain them (Boyan & Sherry, 2011)

Strategies are Mental Models of Games

Claim

Strategies can be treated as mental models of games

- Players develop *mental models* of games to explain them (Boyan & Sherry, 2011)
- Players acquire and *refine* their mental models through practice, instruction and discussion (of strategies) (Boyan et al., 2018)

Strategies are Mental Models of Games

Claim

Strategies can be treated as mental models of games

- Players develop *mental models* of games to explain them (Boyan & Sherry, 2011)
- Players acquire and *refine* their mental models through practice, instruction and discussion (of strategies) (Boyan et al., 2018)
- Strategies can be learned from *other players* (Weintrop & Wilensky, 2013)

Strategies are Mental Models of Games

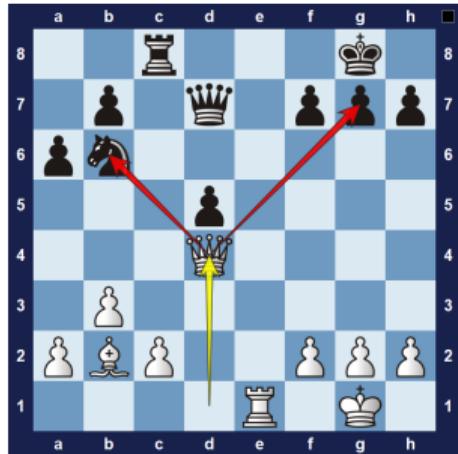
Claim

Strategies can be treated as mental models of games

- Players develop *mental models* of games to explain them (Boyan & Sherry, 2011)
- Players acquire and *refine* their mental models through practice, instruction and discussion (of strategies) (Boyan et al., 2018)
- Strategies can be learned from *other players* (Weintrop & Wilensky, 2013)
- Learned strategies can be *transferred* to one's own gameplay (Paredes-Olay et al., 2002)

Real-world Strategies

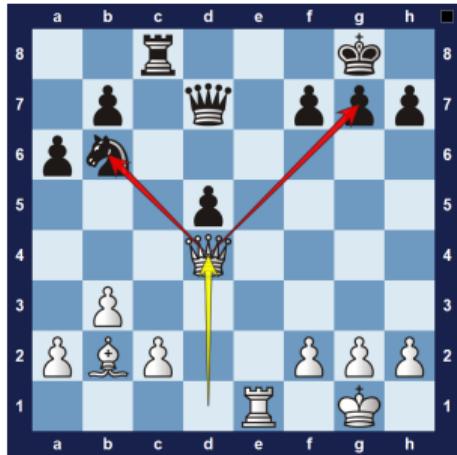
Real-world Strategies



Chessfox

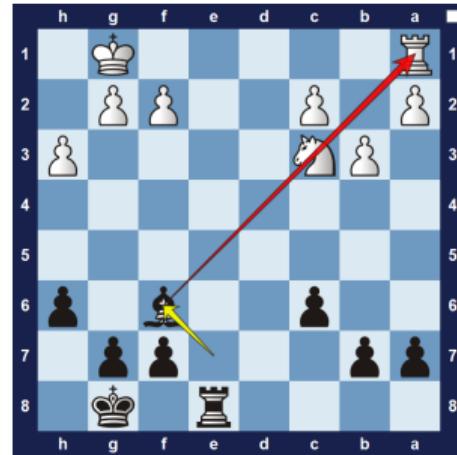
An example of the *fork* tactic in chess

Real-world Strategies



Chessfox

An example of the *fork* tactic in chess



Chessfox

An example of the *pin* tactic in chess

Real-world Strategies



That Blasted Salami

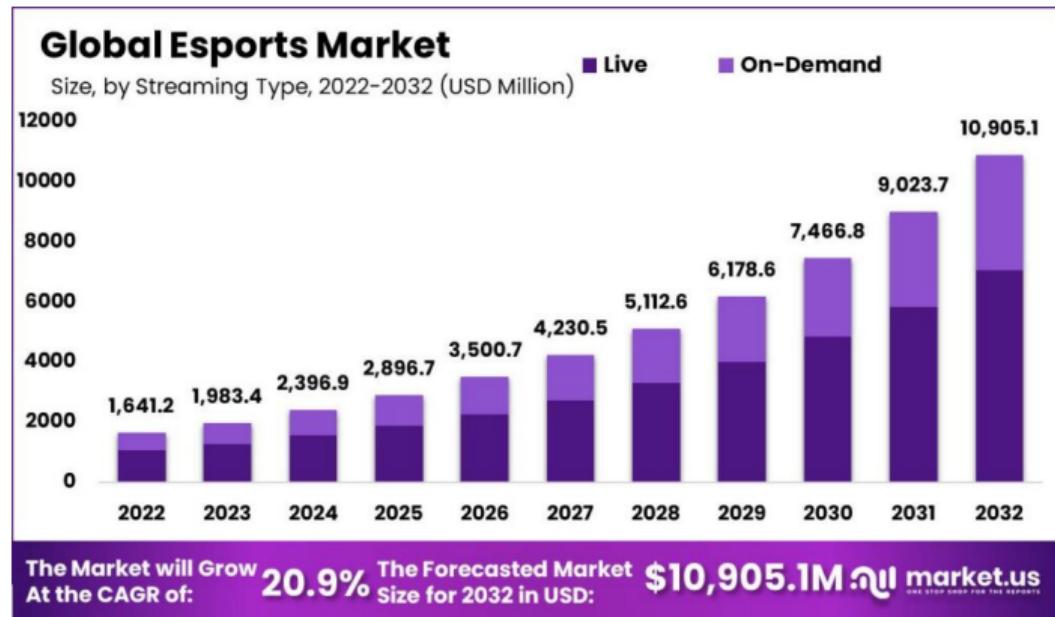
A mixed strategy in *Tekken*

Benefits

Benefits

- Competitive gaming → democratized access to strategies

Esports Viewership



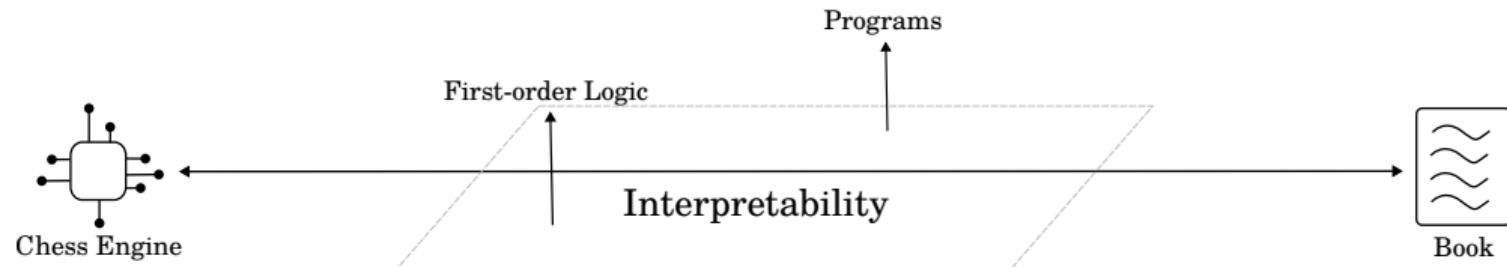
market.us

Benefits

- Competitive gaming → democratized access to strategies
- Explainable AI → techniques for explainability

Benefits

- Competitive gaming → democratized access to strategies
- Explainable AI → techniques for explainability



Thesis Statement

Thesis Statement

A computational model of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of competitive gaming and explainable AI.

RQs

Thesis Statement

A computational model of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of competitive gaming and explainable AI.

RQ1

Thesis Statement

A *computational model* of a game strategy, along with a *synthesis method*, could meet the goals of discovering *high-quality*, *interpretable* strategies and impact the fields of competitive gaming and explainable AI.

RQ1

How do we formally define the problem of *Interpretable Strategy Synthesis* (ISS)?

RQ2

Thesis Statement

A *computational model* of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of *competitive gaming* and explainable AI.

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

RQ3

Thesis Statement

A *computational model* of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of competitive gaming and *explainable AI*.

RQ3

How do we approach the problem of ISS using a *program-based* strategy model?

ISS Framework

RQ1

How do we formally define the problem of *Interpretable Strategy Synthesis* (ISS)?

ISS Framework

RQ1

How do we formally define the problem of *Interpretable Strategy Synthesis* (ISS)?

- Why do we need a framework?

ISS Framework

RQ1

How do we *formally define* the problem of Interpretable Strategy Synthesis (ISS)?

- Why do we need a framework?
- Definitions of key elements —
 - ISS
 - Strategy Model
 - Strategy
 - Performance Measure
 - Interpretability Measure

The Need for a Framework

Paper	Number Used			Interpretability
	Domains	Models	Algorithms	
Spronck, Sprinkhuizen-Kuyper, and Postma (2004)	2	1	1	✗
de Mesentier Silva, Isaksen, Togelius, and Nealen (2016)	1	1	4	✓
Butler, Torlak, and Popović (2017)	1	1	1	✗
Canaan et al. (2018)	1	1	1	✗
de Freitas, de Souza, and Bernardino (2018)	1	1	1	✗
Mariño, Moraes, Oliveira, Toledo, and Lelis (2021)	1	1	1	✗
Krishnan and Martens (2022a)	1	1	1	✗
Mariño and Toledo (2022)	1	1	1	✗
Medeiros, Aleixo, and Lelis (2022)	2	1	2	✗

Table: List of works in ISS

Elements of a Good Framework

Elements of a Good Framework

- Facilitates *comparison*
 - multiple *games*
 - multiple *strategy representations*
 - multiple *learning algorithms*

Elements of a Good Framework

- Facilitates *comparison*
 - multiple *games*
 - multiple *strategy representations*
 - multiple *learning algorithms*
- Provides a clear definition of interpretability

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

- Game environment \mathcal{G}

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

- Game environment \mathcal{G}
- Strategy model \mathcal{M}

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

- Game environment \mathcal{G}
- Strategy model \mathcal{M}
- Performance measure J

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

- Game environment \mathcal{G}
- Strategy model \mathcal{M}
- Performance measure J
- Interpretability measure \mathcal{I}

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Formal

Given a —

- Game environment \mathcal{G}
- Strategy model \mathcal{M}
- Performance measure J
- Interpretability measure \mathcal{I}

The problem of ISS is to find a strategy $\sigma^* \in \mathcal{M}$ that *maximizes both J and \mathcal{I}* .

Strategy Model (\mathcal{M})

- Defines the *space* (or architecture) of strategies

Strategy Model (\mathcal{M})

- Defines the *space* (or architecture) of strategies
- Examples —
 - if-then rules
 - decision trees
 - programmatic scripts
 - neural networks

Strategy (σ)

- A strategy is a *policy* that maps (some) states to actions
- A strategy is an *instance* of a strategy model

Performance Measure (J)

- $J_{\mathcal{G}}(\sigma)$: how *good* a strategy σ is

Performance Measure (J)

- $J_{\mathcal{G}}(\sigma)$: how *good* a strategy σ is
- Understanding the pros/cons of strategies is important

Performance Measure (J)

- $J_{\mathcal{G}}(\sigma)$: how *good* a strategy σ is
- Understanding the pros/cons of strategies is important
- Examples —
 - strategy win rate
 - material advantage (chess)
 - resources harvested (StarCraft II)

Interpretability Measure (\mathcal{I})

- $\mathcal{I}_{\mathcal{G}}(\sigma)$: how *easily understood* a strategy σ is

Interpretability Measure (\mathcal{I})

- $\mathcal{I}_{\mathcal{G}}(\sigma)$: how *easily understood* a strategy σ is
- *computational proxies* based on representation, $\mathcal{I}_{\mathcal{M}}$

Interpretability Measure (\mathcal{I})

- $\mathcal{I}_{\mathcal{G}}(\sigma)$: how *easily understood* a strategy σ is
- *computational proxies* based on representation, $\mathcal{I}_{\mathcal{M}}$
- Examples —
 - number of statements (programmatic script)
 - number of nodes (decision tree)

Interpretability Measure (\mathcal{I})

- $\mathcal{I}_{\mathcal{G}}(\sigma)$: how *easily understood* a strategy σ is
- *computational proxies* based on representation, $\mathcal{I}_{\mathcal{M}}$
- Examples —
 - number of statements (programmatic script)
 - number of nodes (decision tree)
- Use *readable* representations

Summary

RQ1

How do we formally define the problem of Interpretable Strategy Synthesis (ISS)?

Summary

RQ1

How do we *formally define* the problem of Interpretable Strategy Synthesis (ISS)?

- Definition of ISS $\langle \mathcal{G}, \mathcal{M}, J, \mathcal{I} \rangle$ (Krishnan & Martens, 2022b)

Introduction

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

Introduction

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —

Introduction

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —
 - RQ2(a): Applying the ISS framework using FOL-based strategies to chess (Krishnan & Martens, 2022b)

Introduction

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —
 - RQ2(a): Applying the ISS framework using FOL-based strategies to chess (Krishnan & Martens, 2022b)
 - RQ2(b): Learning chess strategies using ILP (Krishnan & Martens, 2022a)

Introduction

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —
 - RQ2(a): Applying the ISS framework using FOL-based strategies to chess (Krishnan & Martens, 2022b)
 - RQ2(b): Learning chess strategies using ILP (Krishnan & Martens, 2022a)
 - RQ2(c): Improving the ILP-based learning method (Krishnan et al., 2024)

First-Order Logic (FOL)

- FOL is a *formal* language

First-Order Logic (FOL)

- FOL is a *formal* language
- can represent *relationships* between objects

First-Order Logic (FOL)

- FOL is a *formal* language
- can represent *relationships* between objects
- E.g., in chess, a king is *in check* if an opponent's piece can capture it

$$\forall k, b \text{ InCheck}(k, b) \iff \exists p \text{ CanMove}(p, \text{Square}(p), \text{Square}(k), b) \wedge p \neq \text{King} \quad (1)$$

Prolog

- *Declarative* programming language

Prolog

- *Declarative* programming language
- Allows specifying *facts* and *rules* in FOL

Prolog

- *Declarative* programming language
- Allows specifying *facts* and *rules* in FOL
- *in check* in Prolog —

```
1 in_check(King , Board) :-  
2     can_move(Piece , Square(Piece) , Square(King) , Board) ,  
3     Piece \= king .
```

Why use FOL-based Strategies?

Why use FOL-based Strategies?

- Prior work in *relational reinforcement learning* and *pattern learning* in chess
 - Handle domains with *large state spaces*
 - *Generalize* to similar domains
 - *Interpretability* of learned rules

Why use FOL-based Strategies?

- Prior work in *relational reinforcement learning* and *pattern learning* in chess
 - Handle domains with *large state spaces*
 - *Generalize* to similar domains
 - *Interpretability* of learned rules
- *Unique capabilities* of rule-based strategy representations

Why use FOL-based Strategies?

- Prior work in *relational reinforcement learning* and *pattern learning* in chess
 - Handle domains with *large state spaces*
 - *Generalize* to similar domains
 - *Interpretability* of learned rules
- *Unique capabilities* of rule-based strategy representations
 - Continuous learning

Why use FOL-based Strategies?

- Prior work in *relational reinforcement learning* and *pattern learning* in chess
 - Handle domains with *large state spaces*
 - *Generalize* to similar domains
 - *Interpretability* of learned rules
- *Unique capabilities* of rule-based strategy representations
 - Continuous learning
 - Predicate invention (Cropper & Dumančić, 2022)

Why use FOL-based Strategies?

- Prior work in *relational reinforcement learning* and *pattern learning* in chess
 - Handle domains with *large state spaces*
 - *Generalize* to similar domains
 - *Interpretability* of learned rules
- *Unique capabilities* of rule-based strategy representations
 - Continuous learning
 - Predicate invention (Cropper & Dumančić, 2022)
- *Interpretability* of rule-based models (Kliegr et al., 2021; Zhang et al., 2021)

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop metrics to show that they suggest better moves than a random baseline?

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a *strategy model* for chess and develop metrics to show that they suggest better moves than a random baseline?

- *Strategy model* for chess

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop *metrics* to show that they suggest better moves than a random baseline?

- Strategy model for chess
- *Performance measure* for chess

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop metrics to *show that they suggest better moves than a random baseline?*

- Strategy model for chess
- Performance measure for chess

Strategy Model for Chess

First-Order (FO) Logic Rule

Strategy Model for Chess

First-Order (FO) Logic Rule

Predicate Vocabulary

Strategy Model for Chess

First-Order (FO) Logic Rule

```
tactic(Position, Move) ←  
    feature_1(…),  
    feature_2(…),  
    :  
    feature_n(…)
```

Predicate Vocabulary

Our chess strategy model expressed in Prolog pseudocode

Strategy Model for Chess

First-Order (FO) Logic Rule

```
tactic(Position, Move) ←  
    feature_1(…),  
    feature_2(…),  
    :  
    feature_n(…)
```

Our chess strategy model expressed in Prolog pseudocode

Predicate Vocabulary

- **Position** =

- [contents(c2,pawn,white),
 contents(g8,knight,black),
 contents(e8,king,black),
 turn(white),kingside_castle(white),...]

Strategy Model for Chess

First-Order (FO) Logic Rule

```
tactic(Position, Move) ←  
    feature_1(…),  
    feature_2(…),  
    :  
    feature_n(…)
```

Our chess strategy model expressed in Prolog pseudocode

Predicate Vocabulary

- **Position** =
[contents(c2,pawn,white),
 contents(g8,knight,black),
 contents(e8,king,black),

 turn(white),kingside_castle(white),...]
- **Move** = [a7,a8,queen]

Strategy Model for Chess

First-Order (FO) Logic Rule

```
tactic(Position, Move) ←  
    feature_1(...),  
    feature_2(...),  
    :  
    feature_n(...)
```

Our chess strategy model expressed in Prolog pseudocode

Predicate Vocabulary

- **Position** = [contents(c2,pawn,white), contents(g8,knight,black), contents(e8,king,black), turn(white), kingside_castle(white), ...]
- **Move** = [a7, a8, queen]
- **Features** =
 - attacks (Pos, Sq1, Sq2)
 - in_check (Pos, Side)
 - is_empty (Pos, Squares)

Example

```
fork(Position,Move) ←  
    legal_move(Position,Move),  
    move(Move,_,To,_),  
    make_move(Position,Move>NewPosition),  
    can_capture(NewPosition,To,ForkSquare1),  
    can_capture(NewPosition,To,ForkSquare2),  
    different(ForkSquare1,ForkSquare2).
```

An interpretation of the *fork* tactic from the chess literature using our predicate vocabulary.

Example

```
fork(Position,Move) ←  
    legal_move(Position,Move),  
    move(Move,_,To,_),  
    make_move(Position,Move,NewPosition),  
    can_capture(NewPosition,To,ForkSquare1),  
    can_capture(NewPosition,To,ForkSquare2),  
    different(ForkSquare1,ForkSquare2).
```

An interpretation of the *fork* tactic from the chess literature using our predicate vocabulary.

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop metrics to show that they suggest better moves than a random baseline?

- Strategy model for chess ✓
- Performance measure for chess

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop *metrics* to show that they suggest better moves than a random baseline?

- Strategy model for chess ✓
- *Performance measure* for chess

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength
- Low divergence → small difference in playing strength

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength
- Low divergence → small difference in playing strength
- Difference in terms of *perceived evaluation* of moves

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength
- Low divergence → small difference in playing strength
- Difference in terms of *perceived evaluation* of moves
- Who is “perceiving”?

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength
- Low divergence → small difference in playing strength
- Difference in terms of *perceived evaluation* of moves
- Who is “perceiving”?
 - Chess-playing agents with an *evaluation function* (chess “engines”)

Performance Measure

Divergence

Equation

- How do strategies *differ* in terms of strength? (Guid & Bratko, 2006, 2011)
- High divergence → large difference in playing strength
- Low divergence → small difference in playing strength
- Difference in terms of *perceived evaluation* of moves
- Who is “perceiving”?
 - Chess-playing agents with an *evaluation function* (chess “engines”)
 - e.g., Stockfish 14 (strong), Maia Chess (McIlroy-Young et al., 2020) (weak, human-like)

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop metrics to show that they suggest better moves than a random baseline?

- Strategy model for chess ✓
- Performance measure for chess ✓

ISS for Chess

RQ2(a)

Could we represent known chess tactics as a strategy model for chess and develop metrics to *show that they suggest better moves than a random baseline?*

- Strategy model for chess ✓
- Performance measure for chess ✓

Evaluation of Known Chess Tactics¹

¹Krishnan and Martens, 2022b.

Evaluation of Known Chess Tactics¹

- PAL (Morales, 1992) $\xrightarrow{\text{learn}}$ *known* chess patterns (tactics) PAL

¹Krishnan and Martens, 2022b.

Evaluation of Known Chess Tactics¹

- PAL (Morales, 1992) $\xrightarrow{\text{learn}}$ known chess patterns (tactics) PAL
- tactics $\xrightarrow{\text{translate}}$ chess strategy model

¹Krishnan and Martens, 2022b.

Evaluation of Known Chess Tactics¹

- PAL (Morales, 1992) $\xrightarrow{\text{learn}}$ known chess patterns (tactics) PAL
- tactics $\xrightarrow{\text{translate}}$ chess strategy model
- Divergence(chess strategies, *human beginner*)

¹Krishnan and Martens, 2022b.

Evaluation of Known Chess Tactics¹

- PAL _(Morales, 1992) $\xrightarrow{\text{learn}}$ known chess patterns (tactics) PAL
- tactics $\xrightarrow{\text{translate}}$ chess strategy model
- Divergence(chess strategies, human beginner)
- Divergence(*random baseline*, human beginner)

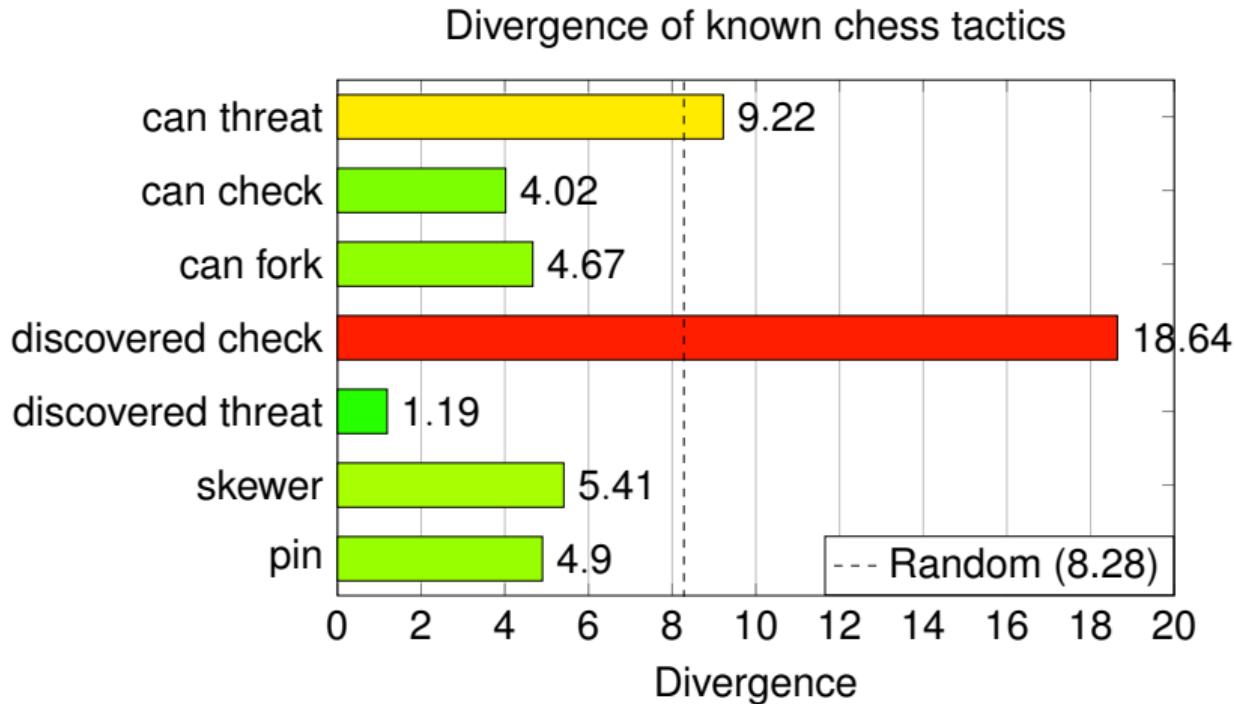
¹Krishnan and Martens, 2022b.

Evaluation of Known Chess Tactics¹

- PAL (Morales, 1992) $\xrightarrow{\text{learn}}$ known chess patterns (tactics) PAL
- tactics $\xrightarrow{\text{translate}}$ chess strategy model
- Divergence(chess strategies, human beginner)
- Divergence(random baseline, human beginner)
- Using a human-like engine (McIlroy-Young et al., 2020)

¹Krishnan and Martens, 2022b.

Results



Learning Chess Strategies

RQ2(b)

Do the chess strategies *learned* using inductive logic programming outperform a random baseline in how closely their divergence scores approximate a beginner player?

- Strategy model for chess ✓
- Performance measure for chess ✓

Learning Chess Strategies

RQ2(b)

Do the chess strategies learned using inductive logic programming outperform a random baseline in how closely their divergence scores approximate a beginner player?

- Strategy model for chess ✓
- Performance measure for chess ✓
- *Learning algorithm* for chess strategies

Learning Chess Strategies

RQ2(b)

Do the chess strategies learned using *inductive logic programming* outperform a random baseline in how closely their divergence scores approximate a beginner player?

- Strategy model for chess ✓
- Performance measure for chess ✓
- *Learning algorithm* for chess strategies

Learning Chess Strategies

RQ2(b)

Do the chess strategies learned using inductive logic programming
outperform a random baseline in how closely their divergence scores approximate a beginner player?

- Strategy model for chess ✓
- Performance measure for chess ✓
- *Learning algorithm* for chess strategies

Inductive Logic Programming (ILP)

Inductive Logic Programming (ILP)

- *symbolic* machine learning technique

Inductive Logic Programming (ILP)

- *symbolic* machine learning technique
- Learn *classifier* in the form of a FOL rule

Inductive Logic Programming (ILP)

- *symbolic* machine learning technique
- Learn *classifier* in the form of a FOL rule
- FOL rule expressed using *background knowledge*

Target Concept

- Want to learn `last` relating a list of characters to the last character in it

$$E^+ = \left\{ \begin{array}{l} \text{last}([m, a, c, h, i, n, e], e). \\ \text{last}([l, e, a, r, n, i, n, g], g). \\ \text{last}([a, l, g, o, r, i, t, h, m], m). \end{array} \right\}$$

$$E^- = \left\{ \begin{array}{l} \text{last}([m, a, c, h, i, n, e], m). \\ \text{last}([m, a, c, h, i, n, e], c). \\ \text{last}([l, e, a, r, n, i, n, g], x). \\ \text{last}([l, e, a, r, n, i, n, g], i). \end{array} \right\}$$

$$B = \left\{ \begin{array}{l} \text{empty}(A) :- \dots \\ \text{head}(A, B) :- \dots \\ \text{tail}(A, B) :- \dots \end{array} \right\}$$

Possible Hypothesis

$$H = \left\{ \begin{array}{l} \text{last}(A, B) :- \text{head}(A, B), \text{tail}(A, C), \text{empty}(C). \\ \text{last}(A, B) :- \text{tail}(A, C), \text{last}(C, B). \end{array} \right\}$$

Learning Rule-based Policies

Learning Rule-based Policies

- Can be learned using ILP

Learning Rule-based Policies

- Can be learned using ILP
- ISS problem $\langle \mathcal{G}, \mathcal{M}, J, \mathcal{I} \rangle \xrightarrow{\text{translate}} \text{ILP problem } \langle E^+, E^-, B \rangle$

Learning Rule-based Policies

- Can be learned using ILP
- ISS problem $\langle \mathcal{G}, \mathcal{M}, J, \mathcal{I} \rangle \xrightarrow{\text{translate}} \text{ILP problem } \langle E^+, E^-, B \rangle$
- E^+ : $\langle s, \pi(s) \rangle$ for some target policy π
- E^- : $\langle s, a' \rangle$ where $a' \in \mathcal{A}(s) - \pi(s)$
- B : manual translation of \mathcal{G} to FOL

Learning Rule-based Policies

- Can be learned using ILP
- ISS problem $\langle \mathcal{G}, \mathcal{M}, J, \mathcal{I} \rangle \xrightarrow{\text{translate}} \text{ILP problem } \langle E^+, E^-, B \rangle$
- E^+ : $\langle s, \pi(s) \rangle$ for some target policy π
- E^- : $\langle s, a' \rangle$ where $a' \in \mathcal{A}(s) - \pi(s)$
- B : manual translation of \mathcal{G} to FOL
- ILP system($\langle E^+, E^-, B \rangle$) $\xrightarrow{\text{learn}}$ rule-based strategies

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system

²Krishnan and Martens, 2022a.

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system
- Modifications —

²Krishnan and Martens, 2022a.

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system
- Modifications —
 - generate only strategies that produce legal moves

²Krishnan and Martens, 2022a.

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system
- Modifications —
 - generate only strategies that produce legal moves
 - prevent generation of strategies that don't match any position in training sets

²Krishnan and Martens, 2022a.

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system
- Modifications —
 - generate only strategies that produce legal moves
 - prevent generation of strategies that don't match any position in training sets
- Evaluate using *divergence* with Stockfish 14 (strong) and Maia Chess (weak, human-like)

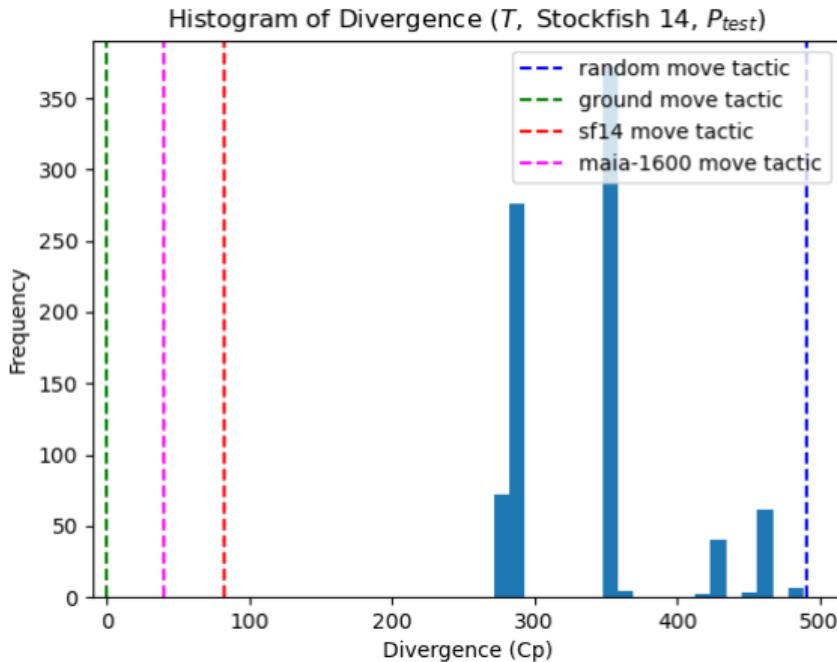
²Krishnan and Martens, 2022a.

Learning Chess Strategies using ILP²

- Use *Popper* (Cropper & Morel, 2021) as ILP system
- Modifications —
 - generate only strategies that produce legal moves
 - prevent generation of strategies that don't match any position in training sets
- Evaluate using *divergence* with Stockfish 14 (strong) and Maia Chess (weak, human-like)
- Compare to random, strong/weak engine baselines

²Krishnan and Martens, 2022a.

Results



Improving the ILP Learning Method

RQ2(c)

Do the chess strategies learned by an ILP system incorporating the changes of the new predicate vocabulary and precision/recall-based constraints produce moves better than those learned by an ILP system without these modifications?

Improving the ILP Learning Method

RQ2(c)

Do the chess strategies learned by an ILP system incorporating the changes of the new predicate vocabulary and *precision/recall-based constraints* produce moves better than those learned by an ILP system without these modifications?

- *Constraints* to limit search space of chess strategies

Improving the ILP Learning Method

RQ2(c)

Do the chess strategies learned by an ILP system incorporating the changes of the new *predicate vocabulary* and precision/recall-based constraints produce moves better than those learned by an ILP system without these modifications?

- Constraints to limit search space of chess strategies
- *New* predicate vocabulary

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies
 - Based on precision and recall of strategy on training set

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies
 - Based on precision and recall of strategy on training set
- New predicate vocabulary — Vocabulary

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies
 - Based on precision and recall of strategy on training set
- New predicate vocabulary — Vocabulary
 - Rules to express *rare* situations – en passant, promotion

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies
 - Based on precision and recall of strategy on training set
- New predicate vocabulary — Vocabulary
 - Rules to express *rare* situations – en passant, promotion
 - Allows *more efficient* search due to introduction of more types

³Krishnan et al., 2024.

Improvements to ILP for Chess Strategies³

- Domain-specific constraints — Constraints
 - *Limit* search space of chess strategies
 - Based on precision and recall of strategy on training set
- New predicate vocabulary — Vocabulary
 - Rules to express *rare* situations – en passant, promotion
 - Allows *more efficient* search due to introduction of more types
 - More domain-specific state features like *can_threat* and *xrays*

³Krishnan et al., 2024.

Evaluation

- Conduct *ablation study* to measure impact of modifications

Evaluation

- Conduct *ablation study* to measure impact of modifications
 - Learn strategies using systems with/without constraints, predicate vocabulary

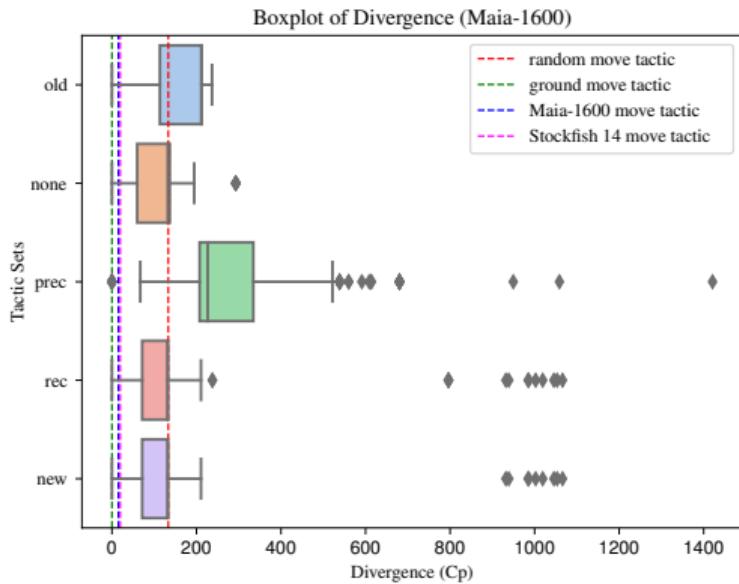
Evaluation

- Conduct *ablation study* to measure impact of modifications
 - Learn strategies using systems with/without constraints, predicate vocabulary
 - Measure average strategy divergence

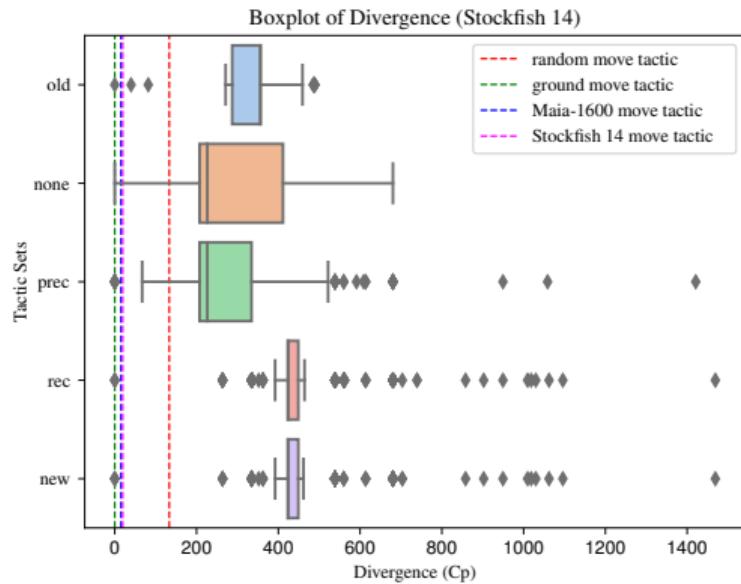
Evaluation

- Conduct *ablation study* to measure impact of modifications
 - Learn strategies using systems with/without constraints, predicate vocabulary
 - Measure average strategy divergence
 - Test decrease vs. old system using a *one-sided Welch's t-test*

Results



Boxplot of strategy divergence (evaluated using *weak* engine) for each system



Boxplot of strategy divergence (evaluated using *strong* engine) for each system

Summary

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

Summary

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —

Summary

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —
 - RQ2(a): Represented chess tactics as FOL rules and showed that known chess tactics resemble human beginner play (Krishnan & Martens, 2022b)

Summary

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —

- RQ2(a): Represented chess tactics as FOL rules and showed that known chess tactics resemble human beginner play (Krishnan & Martens, 2022b)
- RQ2(b): Learned chess strategies using ILP which learned strategies resembling human beginner playing strength (Krishnan & Martens, 2022a)

Summary

RQ2

How do we approach the problem of ISS using a *logic-based* strategy model?

- Application to *chess* —

- RQ2(a): Represented chess tactics as FOL rules and showed that known chess tactics resemble human beginner play (Krishnan & Martens, 2022b)
- RQ2(b): Learned chess strategies using ILP which learned strategies resembling human beginner playing strength (Krishnan & Martens, 2022a)
- RQ2(c): Improving ILP learning using precision/recall constraints (Krishnan et al., 2024)

Challenges with FOL

Challenges with FOL-based strategies —

Challenges with FOL

Challenges with FOL-based strategies —

- Knowledge engineering

Challenges with FOL

Challenges with FOL-based strategies —

- Knowledge engineering
 - Improved predicate vocabulary represented *800+ LoC in Prolog*

Challenges with FOL

Challenges with FOL-based strategies —

- Knowledge engineering
 - Improved predicate vocabulary represented *800+ LoC in Prolog*
- Performance

Challenges with FOL

Challenges with FOL-based strategies —

- Knowledge engineering
 - Improved predicate vocabulary represented *800+ LoC in Prolog*
- Performance
 - Playing strength not close to state-of-the-art

Motivation

Challenges with FOL-based strategies —

- Knowledge engineering
 - Programmatic strategies require *only a DSL* to be defined
- Performance
 - Success of RL methods (DeepMind, 2019; OpenAI et al., 2019; Silver et al., 2018)

Strategies as Programs

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

Strategies as Programs

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

- RQ3(a): Leveraging reinforcement learning (RL) to learn programmatic strategies (Krishnan et al., 2024)

Strategies as Programs

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

- RQ3(a): Leveraging reinforcement learning (RL) to learn programmatic strategies (Krishnan et al., 2024)
- RQ3(b): Showing RL is useful for learning programmatic strategies (Krishnan et al., 2024)

Programmatic Strategy Model

Programmatic Strategy Model

- Programmatic strategy modeled as *domain-specific language (DSL) script*

Programmatic Strategy Model

- Programmatic strategy modeled as *domain-specific language (DSL) script*
- DSL defined as a context-free grammar (CFG)
 $\langle V, \Sigma, R, S \rangle$

$$\begin{aligned}
 P &::= (\epsilon - \mathbf{hd}(h_i)) \\
 I &::= \mathbf{fold}(+, h_i) \\
 D &::= \mathbf{hd}(\mathbf{tl}(h_i)) - \mathbf{hd}(h_i) \\
 C &::= c_1 + c_2 * P + c_3 * I + c_4 * D \\
 B &::= c_0 + c_1 * \mathbf{hd}(h_1) + \dots + c_k * \mathbf{hd}(h_k) > 0 \mid \\
 &\quad B_1 \text{ or } B_2 \mid B_1 \text{ and } B_2 \\
 E &::= C \mid \text{if } B \text{ then } E_1 \text{ else } E_2.
 \end{aligned}$$

The DSL for expressing PID controllers due to Verma et al. (2018).

Programmatic Strategy Model

- Strategy = string in the *language* of the DSL

```
if (0.001 - hd(hTrackPos) > 0) and (0.001 + hd(hTrackPos) > 0)
  then 1.96 + 4.92 * (0.44 - hd(hRPM)) + 0.89 * fold(+, hRPM) + 49.79 * (hd(tl(hRPM)) - hd(hRPM))
  else 1.78 + 4.92 * (0.40 - hd(hRPM)) + 0.89 * fold(+, hRPM) + 49.79 * (hd(tl(hRPM)) - hd(hRPM))
```

A programmatic policy (strategy) for a PID controller.

Learning Programmatic Strategies

- Learning programmatic strategies = program synthesis —

Learning Programmatic Strategies

- Learning programmatic strategies = program synthesis —
 - based on $\langle \text{input}, \text{output} \rangle$ pairs (X. Chen et al., 2019; Lázaro-Gredilla et al., 2019; Yang et al., 2021)

Learning Programmatic Strategies

- Learning programmatic strategies = program synthesis —
 - based on $\langle \text{input}, \text{output} \rangle$ pairs (X. Chen et al., 2019; Lázaro-Gredilla et al., 2019; Yang et al., 2021)
 - demonstrations (Burke et al., 2019; Sun et al., 2018; Xu et al., 2018)

Learning Programmatic Strategies

- Learning programmatic strategies = program synthesis —
 - based on $\langle \text{input}, \text{output} \rangle$ pairs (X. Chen et al., 2019; Lázaro-Gredilla et al., 2019; Yang et al., 2021)
 - demonstrations (Burke et al., 2019; Sun et al., 2018; Xu et al., 2018)

RQ3(a)

How do I formally map the problem of ISS to that of programmatic policy synthesis?

Learning Programmatic Strategies using Decision Transformers⁴

⁴Krishnan et al., 2024.

Learning Programmatic Strategies using Decision Transformers⁴

- ISS for programmatic strategies —

⁴Krishnan et al., 2024.

Learning Programmatic Strategies using Decision Transformers⁴

- ISS for programmatic strategies —
 - Strategy model = DSL

⁴Krishnan et al., 2024.

Learning Programmatic Strategies using Decision Transformers⁴

- ISS for programmatic strategies —
 - Strategy model = DSL
 - Performance measure = reward from executing strategy in target environment

⁴Krishnan et al., 2024.

Learning Programmatic Strategies using Decision Transformers⁴

- ISS for programmatic strategies —
 - Strategy model = DSL
 - Performance measure = reward from executing strategy in target environment
 - Learning method

⁴Krishnan et al., 2024.

Learning Programmatic Strategies using Decision Transformers⁴

- ISS for programmatic strategies —
 - Strategy model = DSL
 - Performance measure = reward from executing strategy in target environment
 - Learning method
- Cast ISS as a RL-based program synthesis problem

⁴Krishnan et al., 2024.

ISS as RL-based Program Synthesis

- Idea: learn a policy to *derive* programs from a CFG

ISS as RL-based Program Synthesis

- Idea: learn a policy to *derive* programs from a CFG

$$S \rightarrow ASB|c$$

$$A \rightarrow aA|\epsilon$$

$$B \rightarrow bB|\epsilon$$

ISS as RL-based Program Synthesis

- Idea: learn a policy to *derive* programs from a CFG

$$S \rightarrow ASB|c$$

$$A \rightarrow aA|\epsilon$$

$$B \rightarrow bB|\epsilon$$

$$S \Rightarrow ASB \Rightarrow SB$$

$$\Rightarrow S \Rightarrow c$$

Program Synthesis MDP

Define a program synthesis MDP \mathcal{M}_{syn} , where —

- States: partially expanded programs
- Actions: rule applications to valid non-terminals
- Transitions: effects of applying a rule
- Rewards: return from executing program in the target MDP

Applying RL to Learn Programmatic Strategies

RQ3b

Could a *decision transformer* model synthesize better strategies for a toy programming environment when compared to existing approaches?

Decision Transformer (DT)⁵

- Application of a transformer model (Vaswani et al., 2017) to RL
- Models the RL problem as a sequence learning problem
- Given an *trajectory* τ , predict the next action

$$\tau = (\hat{R}_1, s_1, a_1, \hat{R}_2, s_2, a_2, \dots, \hat{R}_T, s_T, a_T) \quad (2)$$

⁵Y. Chen, Wang, Bastani, Dillig, and Feng, 2020.

Methodology

Modifications to original decision transformer (DT) architecture —

Methodology

Modifications to original decision transformer (DT) architecture —

- *Action Masking* to deal with invalid actions (Huang & Ontañón, 2022)

Methodology

Modifications to original decision transformer (DT) architecture —

- *Action Masking* to deal with invalid actions (Huang & Ontañón, 2022)
- *Cross-entropy loss* and *sampling from action distribution* to deal with discrete action spaces

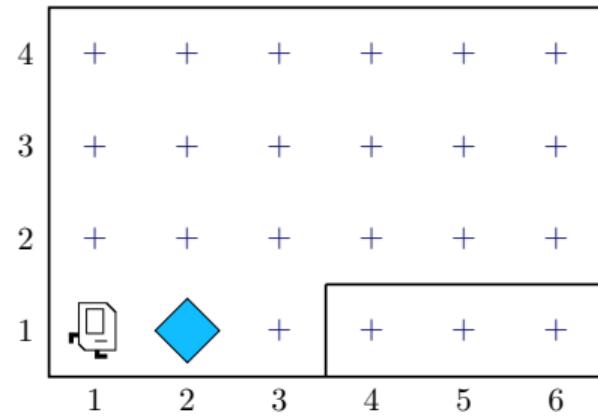
Methodology

Modifications to original decision transformer (DT) architecture —

- *Action Masking* to deal with invalid actions (Huang & Ontañón, 2022)
- *Cross-entropy loss* and *sampling from action distribution* to deal with discrete action spaces
- *GRU-based state embedding* to better encode state information

Karel Domain

- Simple programming language and environment for teaching
- Discrete action space
- Use task definitions from Trivedi, Zhang, Sun, and Lim (2021)



A sample Karel world

Training & Evaluation

- Randomly generated program dataset

Training & Evaluation

- Randomly generated program dataset
- Programs $\xrightarrow{\text{convert}}$ trajectories in \mathcal{M}_{syn}

Training & Evaluation

- Randomly generated program dataset
- Programs $\xrightarrow{\text{convert}}$ trajectories in \mathcal{M}_{syn}
- Train modified DT on generated trajectories

Training & Evaluation

- Randomly generated program dataset
- Programs $\xrightarrow{\text{convert}}$ trajectories in \mathcal{M}_{syn}
- Train modified DT on generated trajectories
- Parameterize top- $k\%$ of the programs based on reward to train on

Training & Evaluation

- Randomly generated program dataset
- Programs $\xrightarrow{\text{convert}}$ trajectories in \mathcal{M}_{syn}
- Train modified DT on generated trajectories
- Parameterize top- $k\%$ of the programs based on reward to train on
- Compare performance with LEAPS (Trivedi et al., 2021)

Sample Program

```
1 DEF run m(
2     WHILE c( noMarkersPresent c) w(
3         turnRight
4         move
5     w)
6     IF c( rightIsClear c) i(
7         WHILE c( frontIsClear c) w(
8             move
9         w)
0     i )
1 m)
```

A program learned by the DT

Program Description

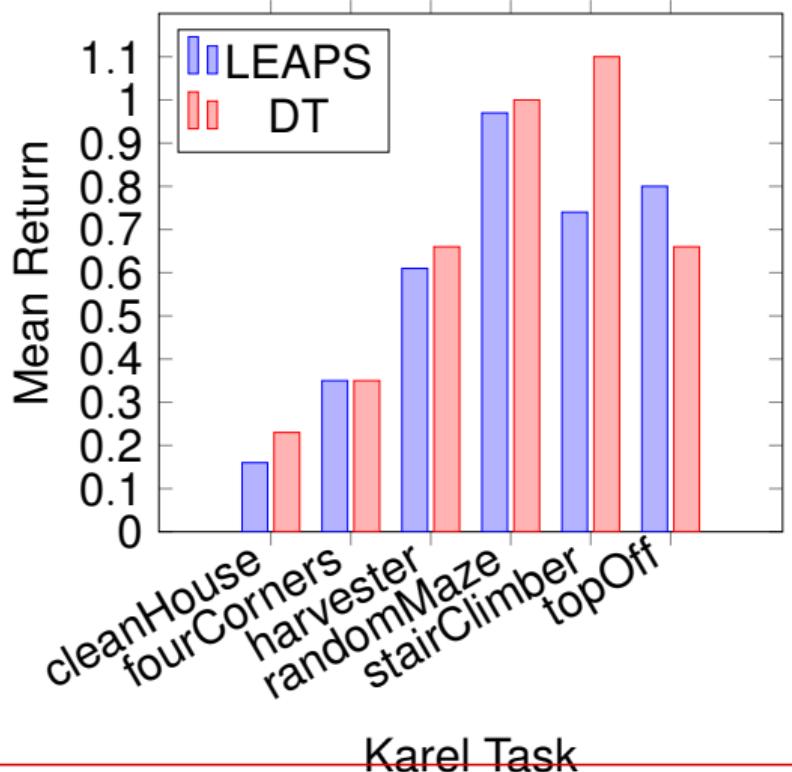
Strategy description as generated by Claude 3.5 Sonnet (Anthropic, 2024)

Claude 3.5 Sonnet Description

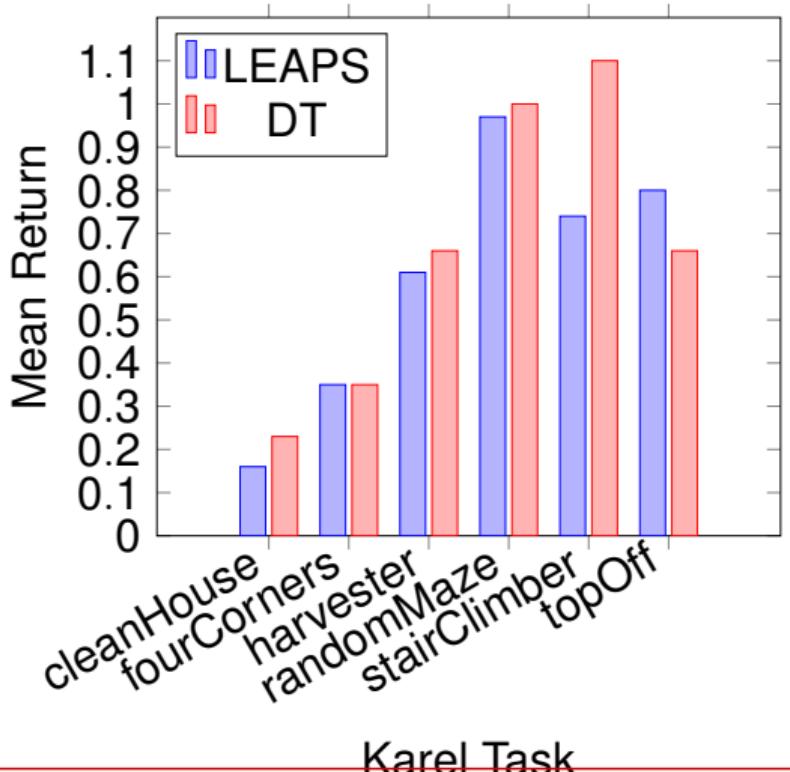
In summary, the strategy is:

- ➊ Climb the stairs by repeatedly turning right and moving forward until a marker is found.
- ➋ Once at the top (indicated by the marker), if there's clear space to the right, move forward along the top of the stairs until reaching an obstacle.

Results



Results



Task	Unique Programs	
	LEAPS	DT
cleanHouse	3627	59
fourCorners	9872	55
harvester	11708	28
randomMaze	295	63
stairClimber	298	49
topOff	30278	63

Table: Number of unique programs explored by LEAPS and the DT during their search (sampling) phases.

Sample Efficiency

- Search
 - LEAPS uses CEM search
 - DT uses random sampling
 - DT explores orders of magnitude fewer programs

Sample Efficiency

- Search
 - LEAPS uses CEM search
 - DT uses random sampling
 - DT explores orders of magnitude fewer programs
- Training
 - Optimal top- $k\%$ = 0.1%, implies DT trained on 500 programs
 - LEAPS trained on 35,000 programs

Sample Efficiency

- Search
 - LEAPS uses CEM search
 - DT uses random sampling
 - DT explores orders of magnitude fewer programs
- Training
 - Optimal $\text{top-}k\% = 0.1\%$, implies DT trained on 500 programs
 - LEAPS trained on 35,000 programs
- Both achieve comparable performance on almost all tasks

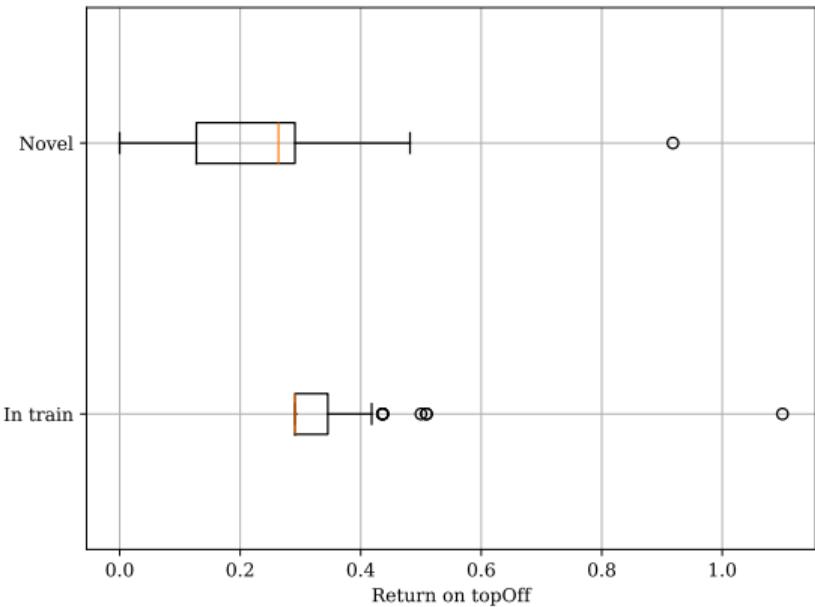
Novelty of Generated Programs

Novelty of Generated Programs

- Plotted distribution of rewards for novel and non-novel programs

Novelty of Generated Programs

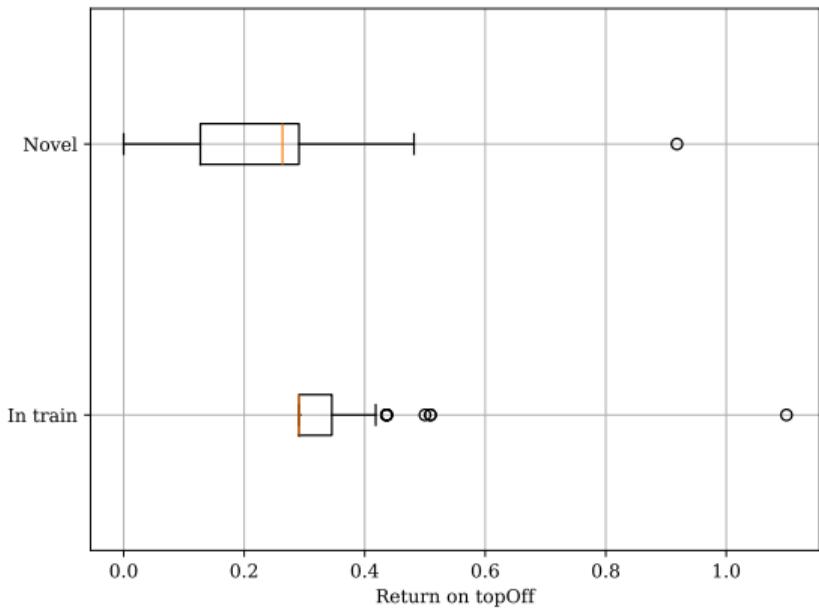
- Plotted distribution of rewards for novel and non-novel programs



Box plot of returns for programs sampled from the decision transformer model.

Novelty of Generated Programs

- Plotted distribution of rewards for novel and non-novel programs
- Mean reward for novel programs is *lower* than non-novel programs ($p < 0.01$)



Box plot of returns for programs sampled from the decision transformer model.

Summary

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

Summary

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

- RQ3(a): Program synthesis MDP (Krishnan et al., 2024)

Summary

RQ3

How do we approach the problem of ISS using a *programmatic* strategy model?

- RQ3(a): Program synthesis MDP (Krishnan et al., 2024)
- RQ3(b): Decision transformer matches SOTA performance with less training data and search (Krishnan et al., 2024)

Conclusion

Thesis Statement

A computational model of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of competitive gaming and explainable AI.

Contributions —

Conclusion

Thesis Statement

A *computational model* of a game strategy, along with a synthesis method, could meet the goals of discovering *high-quality*, *interpretable* strategies and impact the fields of competitive gaming and explainable AI.

Contributions —

- RQ1: Defined a *framework* for ISS (Krishnan & Martens, 2022b)

Conclusion

Thesis Statement

A computational model of a game strategy, along with a synthesis method, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of *competitive gaming* and explainable AI.

Contributions —

- RQ2: Applied *FOL-based* strategy model to solve ISS for chess

Conclusion

Thesis Statement

A *computational model* of a game strategy, along with a synthesis method, could meet the goals of discovering *high-quality*, *interpretable* strategies and impact the fields of *competitive gaming* and explainable AI.

Contributions —

- RQ2: Applied *FOL-based* strategy model to solve ISS for chess
 - RQ2(a): Represented chess tactics as FOL rules (Krishnan & Martens, 2022b)

Conclusion

Thesis Statement

A computational model of a game strategy, along with a *synthesis method*, could meet the goals of discovering high-quality, interpretable strategies and impact the fields of *competitive gaming* and explainable AI.

Contributions —

- RQ2: Applied *FOL-based* strategy model to solve ISS for chess
 - RQ2(a): Represented chess tactics as FOL rules (Krishnan & Martens, 2022b)
 - RQ2(b): Learned chess strategies using ILP (Krishnan & Martens, 2022a)

Conclusion

Thesis Statement

A computational model of a game strategy, along with a *synthesis method*, could meet the goals of discovering *high-quality*, interpretable strategies and impact the fields of *competitive gaming* and explainable AI.

Contributions —

- RQ2: Applied *FOL-based* strategy model to solve ISS for chess
 - RQ2(a): Represented chess tactics as FOL rules (Krishnan & Martens, 2022b)
 - RQ2(b): Learned chess strategies using ILP (Krishnan & Martens, 2022a)
 - RQ2(c): Improving ILP learning (Krishnan et al., 2024)

Conclusion

Thesis Statement

A *computational model* of a game strategy, along with a *synthesis method*, could meet the goals of discovering *high-quality*, interpretable strategies and impact the fields of competitive gaming and *explainable AI*.

Contributions —

- RQ3: Applied *programmatic* strategy model to solve ISS using RL

Conclusion

Thesis Statement

A computational model of a game strategy, along with a *synthesis method*, could meet the goals of discovering *high-quality*, *interpretable* strategies and impact the fields of competitive gaming and *explainable AI*.

Contributions —

- RQ3: Applied *programmatic* strategy model to solve ISS using RL
 - RQ3(a): Program synthesis MDP (Krishnan et al., 2024)

Conclusion

Thesis Statement

A computational model of a game strategy, along with a *synthesis method*, could meet the goals of discovering *high-quality*, interpretable strategies and impact the fields of competitive gaming and *explainable AI*.

Contributions —

- RQ3: Applied *programmatic* strategy model to solve ISS using RL
 - RQ3(a): Program synthesis MDP (Krishnan et al., 2024)
 - RQ3(b): Decision transformer matches SOTA performance with less training data and search (Krishnan et al., 2024)

Publications

- Krishnan, Abhijeet, Colin M. Potts, Arnav Jhala, Harshad Khadilkar, Shirish Karande and Chris Martens. "Learning Explainable Representations of Complex Game-playing Strategies." *Proceedings of the Eleventh Annual Conference on Advances in Cognitive Systems*. 2024.
- Villalobos-Arias, Leonardo, Derek Martin, Abhijeet Krishnan, Madeleine Gagné, Colin M. Potts and Arnav Jhala. "Modeling Risk in Reinforcement Learning: A Literature Mapping." *arXiv preprint arXiv:2312.05231*. 2023.
- Krishnan, Abhijeet and Chris Martens. "Synthesizing Chess Tactics from Player Games." In *Workshop on Artificial Intelligence for Strategy Games (SG) and Esports Analytics (EA), 18th AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*. 2022.
- Krishnan, Abhijeet and Chris Martens. "Towards the Automatic Synthesis of Interpretable Chess Tactics." In *Explainable Agency in Artificial Intelligence Workshop, 36th AAAI Conference on Artificial Intelligence*. 2022.
- Krishnan, Abhijeet, Aaron Williams, and Chris Martens. "Towards Action Model Learning for Player Modeling." *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*. Vol. 16. No. 1. 2020.
- Krishnan, Abhijeet and Chris Martens. "Rule-based Cognitive Modeling via Human-Computer Interaction." Poster presented at: *5th LAS Research Symposium*; 2019 Dec 10; Raleigh, NC.

Thank You!

Questions?

References I

- Anthropic. (2024). The claude 3.5 sonnet model card addendum.
Retrieved from https://www-cdn.anthropic.com/fed9cc193a14b84131812372d8d5857f8f304c52/Model_Card_Claude_3_Addendum.pdf
- Beau, P., & Bakkes, S. (2016, September). Automated game balancing of asymmetric video games. In *2016 ieee conference on computational intelligence and games (cig)* (pp. 1–8). doi:10.1109/CIG.2016.7860432

References II

- Bergdahl, J., Gordillo, C., Tollmar, K., & Gisslén, L. (2020, August). Augmenting automated game testing with deep reinforcement learning. In *2020 ieee conference on games (cog)* (pp. 600–603). doi:10.1109/CoG47356.2020.9231552
- Berliner, H. J. (1975). *A representation and some mechanisms for a problem solving chess program*. Carnegie-Mellon Univ Pittsburgh PA Dept of Computer Science.

References III

- Boyan, A., McGloin, R., & Wasserman, J. A. (2018). Model matching theory: A framework for examining the alignment between game mechanics and mental models. *Media and Communication*, 6(2), 126–136.
doi:<https://doi.org/10.17645/mac.v6i2.1326>
- Boyan, A., & Sherry, J. L. (2011). The challenge in creating games for education: Aligning mental models with game models. *Child Development Perspectives*, 5(2), 82–87.
doi:10.1111/j.1750-8606.2011.00160.x

References IV

- Bramer, M. A. (1977). *Representation of knowledge for chess endgames towards a self-improving system.* (Doctoral dissertation, Open University (United Kingdom)).
- Bratko, I. (1982). Knowledge-based problem-solving in al3. *Machine intelligence*, 10, 73–100.
- Burke, M., Penkov, S. V., & Ramamoorthy, S. (2019, June). From explanation to synthesis: Compositional program induction for learning from demonstration. In *Robotics: Science and systems XV*. doi:10.15607/rss.2019.xv.015

References V

- Butler, E., Torlak, E., & Popović, Z. (2017). Synthesizing interpretable strategies for solving puzzle games. In *Proceedings of the 12th international conference on the foundations of digital games*. doi:10.1145/3102071.3102084
- Canaan, R., Shen, H., Torrado, R., Togelius, J., Nealen, A., & Menzel, S. (2018). Evolving agents for the hanabi 2018 cig competition. In *2018 ieee conference on computational intelligence and games (cig)* (pp. 1–8). doi:10.1109/CIG.2018.8490449

References VI

- Chen, X., Liu, C., & Song, D. (2019). Execution-guided neural program synthesis. In *International conference on learning representations*. Retrieved from <https://openreview.net/forum?id=H1gfOiAqYm>
- Chen, Y., Wang, C., Bastani, O., Dillig, I., & Feng, Y. (2020). Program Synthesis Using Deduction-Guided Reinforcement Learning. In S. K. Lahiri & C. Wang (Eds.), *Computer Aided Verification* (pp. 587–610). doi:10.1007/978-3-030-53291-8_30

References VII

- Chitayat, A. P., Block, F., Walker, J., & Drachen, A. (2023). Beyond the meta: Leveraging game design parameters for patch-agnostic esport analitics. *Proceedings of the AAAI Conference on Artificial Intelligence and Interactive Digital Entertainment*, 19(11), 116–125. doi:10.1609/aiide.v19i1.27507
- Clark, N., Macdonald, B., & Kloo, I. (2020). A bayesian adjusted plus-minus analysis for the esport dota 2. *Journal of Quantitative Analysis in Sports*, 16(4), 325–341. doi:10.1515/jqas-2019-0103

References VIII

- Cropper, A., & Dumančić, S. (2022). Inductive logic programming at 30: A new introduction. *Journal of Artificial Intelligence Research*, 74, 765–850. doi:10.1613/jair.1.13507
- Cropper, A., & Morel, R. (2021). Learning programs by learning from failures. *Machine Learning*, 110(4), 801–856.
- DeepMind. (2019, January). Alphastar: Mastering the real-time strategy game starcraft ii. Retrieved September 21, 2022, from <https://www.deepmind.com/blog/alphastar-mastering-the-real-time-strategy-game-starcraft-ii>

References IX

- de Freitas, J. M., de Souza, F. R., & Bernardino, H. S. (2018). Evolving controllers for mario ai using grammar-based genetic programming. In *2018 ieee congress on evolutionary computation (cec)* (pp. 1–8). doi:10.1109/CEC.2018.8477698
- DeLaurentis, D. A., Panchal, J. H., Raz, A. K., Balasubramani, P., Maheshwari, A., Dachowicz, A., & Mall, K. (2021, August). Toward automated game balance: A systematic engineering design approach. In *2021 ieee conference on games (cog)* (pp. 1–8). doi:10.1109/CoG52621.2021.9619032

References X

- de Mesentier Silva, F., Isaksen, A., Togelius, J., & Nealen, A. (2016). Generating heuristics for novice players. In *2016 ieee conference on computational intelligence and games (cig)* (pp. 1–8). IEEE.
- Evans, R., & Grefenstette, E. (2018). Learning explanatory rules from noisy data. *Journal of Artificial Intelligence Research*, *61*, 1–64.

References XI

Frommel, J., Gugenheimer, J., & Rogers, K. (2019). Opportunities and challenges of using game video stream data for games research. In *All the (world wide) webs a stage: A workshop on live streaming, workshop at chi 2019.*

Gebser, M., Harrison, A., Kaminski, R., Lifschitz, V., & Schaub, T. (2015). Abstract gringo. *Theory and Practice of Logic Programming*, 15(4-5), 449–463.

References XII

- Gelfond, M., & Lifschitz, V. (1988). The stable model semantics for logic programming.. In *Iclp/slp* (Vol. 88, pp. 1070–1080). Cambridge, MA.
- Gobet, F., & Jansen, P. J. (2006). Training in chess: A scientific approach. *Education and chess*.
- Guid, M., & Bratko, I. (2006). Computer analysis of world chess champions. *ICGA journal*, 29(2), 65–73.

References XIII

- Guid, M., & Bratko, I. (2011). Using heuristic-search based engines for estimating human skill at chess. *ICGA journal*, 34(2), 71–81.
- Hodge, V. J., Devlin, S., Sephton, N., Block, F., Cowling, P. I., & Drachen, A. (2021). Win prediction in multiplayer esports: Live professional match prediction. *IEEE Transactions on Games*, 13(4), 368–379. doi:10.1109/TG.2019.2948469

References XIV

- Huang, S., & Ontañón, S. (2022). A Closer Look at Invalid Action Masking in Policy Gradient Algorithms. *The International FLAIRS Conference Proceedings, 35*. arXiv:2006.14171 [cs, stat]. doi:10.32473/flairs.v35i.130584
- Huberman, B. J. (1968, July). *A program to play chess end games.* (Doctoral dissertation, Department of Computer Science, Stanford University.).

References XV

- Jennings-Teats, M., Smith, G., & Wardrip-Fruin, N. (2010). Polymorph: Dynamic difficulty adjustment through level generation. In *Proceedings of the 2010 workshop on procedural content generation in games* (pp. 1–4).
- Kliegr, T., Bahník, Š., & Fürnkranz, J. (2021). A review of possible effects of cognitive biases on interpretation of rule-based machine learning models. *Artificial Intelligence*, 295, 103458. doi:<https://doi.org/10.1016/j.artint.2021.103458>

References XVI

Krishnan, A., & Martens, C. (2022a, October). Synthesizing interpretable chess tactics from player games. In *Proceedings of the workshop on artificial intelligence for strategy games (sg) and esports analytics (ea), 18th aaai conference on artificial intelligence and interactive digital entertainment*, American Association for Artificial Intelligence.

References XVII

Krishnan, A., & Martens, C. (2022b, March). Towards the automatic synthesis of interpretable chess tactics. In *Proceedings of the explainable agency in artificial intelligence workshop, 36th aaai conference on artificial intelligence* (pp. 91–97). American Association of Artificial Intelligence.

References XVIII

- Krishnan, A., Potts, C. M., Jhala, A., Khadilkar, H., Karande, S., & Martens, C. (2024, June). Learning explainable representations of complex game-playing strategies. In *Proceedings of the eleventh annual conference on advances in cognitive systems*. (to appear).
- Lázaro-Gredilla, M., Lin, D., Guntupalli, J. S., & George, D. (2019). Beyond imitation: Zero-shot task transfer on robots by learning concepts as cognitive programs. *Science Robotics*, 4(26), eaav3150. doi:10.1126/scirobotics.aav3150

References XIX

- Mariño, J. R. H., Moraes, R. O., Oliveira, T. C., Toledo, C., & Lelis, L. H. S. (2021). Programmatic strategies for real-time strategy games. *Proceedings of the AAAI Conference on Artificial Intelligence*, 35(1), 381–389.
doi:10.1609/aaai.v35i1.16114
- Mariño, J. R., & Toledo, C. F. (2022). Evolving interpretable strategies for zero-sum games. *Applied Soft Computing*, 122, 108860.

References XX

- Maymin, P. Z. (2021). Smart kills and worthless deaths: Esports analytics for league of legends. *Journal of Quantitative Analysis in Sports*, 17(1), 11–27. doi:10.1515/jqas-2019-0096
- McIlroy-Young, R., Sen, S., Kleinberg, J., & Anderson, A. (2020). Aligning superhuman ai with human behavior: Chess as a model system. In *Proceedings of the 26th acm sigkdd international conference on knowledge discovery & data mining* (pp. 1677–1687). doi:10.1145/3394486.3403219

References XXI

Medeiros, L. C., Aleixo, D. S., & Lelis, L. H. S. (2022). What can we learn even from the weakest? learning sketches for programmatic strategies. (arXiv:2203.11912).

arXiv:2203.11912 [cs]. Retrieved from
<http://arxiv.org/abs/2203.11912>

Morales, E. (1992). *First order induction of patterns in chess*. (Doctoral dissertation, PhD thesis, The Turing Institute-University of Strathclyde).

References XXII

Mousavinasab, E., Zarifsanaiey, N., R. Niakan Kalhor, S., Rakhshan, M., Keikha, L., & Ghazi Saeedi, M. (2021). Intelligent tutoring systems: A systematic review of characteristics, applications, and evaluation methods. *Interactive Learning Environments*, 29(1), 142–163. doi:10.1080/10494820.2018.1558257

References XXIII

- Nagel, G. L. (2017, January). *Use of eye tracking for esports analytics in a moba game*. (Master thesis, The University of Bergen). Accepted: 2017-05-18T12:37:39Z. Retrieved from <https://bora.uib.no/bora-xmlui/handle/1956/15875>
- Novak, A. R., Bennett, K. J., Pluss, M. A., & Fransen, J. (2020). Performance analysis in esports: Modelling performance at the 2018 league of legends world championship. *International Journal of Sports Science & Coaching*, 15(56), 809–817. doi:10.1177/1747954120932853

References XXIV

- OpenAI, : Berner, C., Brockman, G., Chan, B., Cheung, V., ...
Zhang, S. (2019). Dota 2 with large scale deep reinforcement learning. arXiv: 1912.06680 [cs.LG]
- Paredes-Olay, C., Abad, M. J. F., Gámez, M., & Rosas, J. M. (2002). Transfer of control between causal predictive judgments and instrumental responding. *Animal Learning & Behavior*, 30(3), 239–248. doi:10.3758/BF03192833

References XXV

Pfau, J., Smeddinck, J. D., & Malaka, R. (2017). Automated game testing with icarus: Intelligent completion of adventure riddles via unsupervised solving. In *Extended abstracts publication of the annual symposium on computer-human interaction in play* (pp. 153–164). doi:10.1145/3130859.3131439

Pitrat, J. (1977). A chess combination program which uses plans. *Artificial Intelligence*, 8(3), 275–321.

References XXVI

- Qiu, W., & Zhu, H. (2022). Programmatic reinforcement learning without oracles. In *International conference on learning representations*. Retrieved from <https://openreview.net/forum?id=6Tk2noBdvxt>
- Schubert, M., Drachen, A., & Mahlmann, T. (2016). Esports analytics through encounter detection. In *Proceedings of the mit sloan sports analytics conference* (Vol. 1, p. 2016). MIT Sloan Boston, MA.

References XXVII

Silver, D., Hubert, T., Schrittwieser, J., Antonoglou, I., Lai, M., Guez, A., ... Hassabis, D. (2018). A general reinforcement learning algorithm that masters chess, shogi, and Go through self-play. *Science*, 362(6419), 1140–1144. Publisher: American Association for the Advancement of Science. doi:10.1126/science.aar6404

References XXVIII

- Smith, A. M., Lewis, C., Hullet, K., Smith, G., & Sullivan, A. (2011). An inclusive view of player modeling. In *Proceedings of the 6th international conference on foundations of digital games* (pp. 301–303).
- Spronck, P., Sprinkhuizen-Kuyper, I., & Postma, E. (2004). Online adaptation of game opponent ai with dynamic scripting. *International Journal of Intelligent Games and Simulation*, 3(1), 45–53.

References XXIX

Stepanov, A., Lange, A., Khromov, N., Korotin, A., Burnaev, E., & Somov, A. (2019, July). Sensors and game synchronization for data analysis in esports. In *2019 ieee 17th international conference on industrial informatics (indin)* (Vol. 1, pp. 933–938). doi:10.1109/INDIN41052.2019.8972249

References XXX

- Sun, S.-H., Noh, H., Somasundaram, S., & Lim, J. (2018, July). Neural program synthesis from diverse demonstration videos. In J. Dy & A. Krause (Eds.), *Proceedings of the 35th international conference on machine learning* (Vol. 80, pp. 4790–4799). PMLR. Retrieved from <https://proceedings.mlr.press/v80/sun18a.html>
- Sutton, R. S., & Barto, A. G. (2018). *Reinforcement learning: An introduction*. MIT press.

References XXXI

Szabo, A. (1984). *Computer chess tactics and strategy*.
(Doctoral dissertation, University of British Columbia).
doi:<http://dx.doi.org/10.14288/1.0051870>

References XXXII

Trivedi, D., Zhang, J., Sun, S.-H., & Lim, J. J. (2021). Learning to Synthesize Programs as Interpretable and Generalizable Policies. In *Advances in Neural Information Processing Systems* (Vol. 34, pp. 25146–25163). Curran Associates, Inc. Retrieved June 7, 2023, from <https://proceedings.neurips.cc/paper/2021/hash/d37124c4c79f357cb02c655671a432fa-Abstract.html>

References XXXIII

Vaswani, A., Shazeer, N., Parmar, N., Uszkoreit, J., Jones, L., Gomez, A. N., ... Polosukhin, I. (2017). Attention is all you need. In I. Guyon, U. V. Luxburg, S. Bengio, H. Wallach, R. Fergus, S. Vishwanathan, & R. Garnett (Eds.), *Advances in neural information processing systems* (Vol. 30), Curran Associates, Inc. Retrieved from https://proceedings.neurips.cc/paper_files/paper/2017/file/3f5ee243547dee91fb053c1c4a845aa-Paper.pdf

References XXXIV

Verma, A., Murali, V., Singh, R., Kohli, P., & Chaudhuri, S. (2018, July). Programmatically interpretable reinforcement learning. In J. Dy & A. Krause (Eds.), *Proceedings of the 35th international conference on machine learning* (Vol. 80, pp. 5045–5054). PMLR. Retrieved from <https://proceedings.mlr.press/v80/verma18a.html>

References XXXV

- Weintrop, D., & Wilensky, U. (2013). Know your enemy: Learning from in-game opponents. In *Proceedings of the 12th international conference on interaction design and children* (pp. 408–411). doi:10.1145/2485760.2485789
- Wilkins, D. E. (1979). *Using patterns and plans to solve problems and control search*. Stanford University.
- Xenopoulos, P., Rulff, J., & Silva, C. (2022). Ggviz: Accelerating large-scale esports game analysis. *Proc. ACM Hum.-Comput. Interact.*, 6(CHI PLAY), 238:1–238:22. doi:10.1145/3549501

References XXXVI

Xu, D., Nair, S., Zhu, Y., Gao, J., Garg, A., Fei-Fei, L., & Savarese, S. (2018). Neural task programming: Learning to generalize across hierarchical tasks. In *2018 ieee international conference on robotics and automation (icra)* (pp. 3795–3802). doi:10.1109/ICRA.2018.8460689

References XXXVII

- Xue, S., Wu, M., Kolen, J., Aghdaie, N., & Zaman, K. A. (2017). Dynamic difficulty adjustment for maximized engagement in digital games. In *Proceedings of the 26th international conference on world wide web companion* (pp. 465–471). doi:10.1145/3041021.3054170

References XXXVIII

- Yang, Y. D., Inala, J. P., Bastani, O., Pu, Y., Solar-Lezama, A., & Rinard, M. (2021, November). Program Synthesis Guided Reinforcement Learning for Partially Observed Environments. arXiv:2102.11137 [cs]. arXiv. Retrieved June 28, 2023, from <http://arxiv.org/abs/2102.11137>
- Zhang, Y., Tio, P., Leonardis, A., & Tang, K. (2021). A survey on neural network interpretability. *IEEE Transactions on Emerging Topics in Computational Intelligence*, 5(5), 726–742. doi:10.1109/TETCI.2021.3100641

References XXXIX

Zheng, Y., Xie, X., Su, T., Ma, L., Hao, J., Meng, Z., ... Fan, C. (2019, November). Wuji: Automatic online combat game testing using evolutionary deep reinforcement learning. In *2019 34th ieee/acm international conference on automated software engineering (ase)* (pp. 772–784). doi:10.1109/ASE.2019.00077

Interpretable Strategy Synthesis (ISS)

Definition (ISS)

Given a —

- Game environment \mathcal{G} (i.e., an MDP $\langle \mathcal{S}, \mathcal{A}(s), \mathcal{P}, \mathcal{R}, \gamma \rangle$)
- Strategy model \mathcal{M}
- Performance measure $J: \mathcal{M} \times \mathcal{G} \rightarrow \mathbb{R}$
- Interpretability measure $\mathcal{I}: \mathcal{M} \times \mathcal{G} \rightarrow \mathbb{R}$

The problem of ISS is to find a strategy σ^* s.t. —

$$\sigma^* \doteq \arg \max_{\sigma} J(\sigma, \mathcal{G}) \mathcal{I}(\sigma, \mathcal{G}), \sigma \in \mathcal{M}$$

Strategy Model (\mathcal{M})

Definition (strategy model)

A strategy model (\mathcal{M}) is a function $\mathcal{M}: \mathcal{S} \times \Theta \rightarrow \mathcal{A}$, where —

- \mathcal{S} is the state space of \mathcal{G} ,
- Θ is the associated parameter space of \mathcal{M} , and
- \mathcal{A} is the action space of \mathcal{G}

◀ Return

Strategy (σ)

Definition (strategy)

A strategy (σ) for \mathcal{G} is a probability distribution over the available actions in a state, for a *subset* of states in the state space (\mathcal{S}), and parameterized by $\theta \in \Theta$. The states for which σ is defined are termed as the applicable states set, and are given by $A_\sigma \subseteq \mathcal{S}$.

$$\sigma: \mathcal{S} \times \mathcal{A} \times \Theta \rightarrow [0, 1] \quad (3)$$

[◀ Return](#)

Performance Measure (J)

Definition (performance measure)

The performance measure (J) is a function $J: \mathcal{M} \times \mathcal{G} \rightarrow \mathbb{R}$ that evaluates the performance of a strategy $\sigma \in \mathcal{M}$, in a game environment \mathcal{G} , by assigning it a numerical score.

◀ Return

Interpretability Measure (\mathcal{I})

Definition (interpretability measure)

Formal

The interpretability measure (\mathcal{I}) is a function $\mathcal{I}: \mathcal{M} \times \mathcal{G} \rightarrow \mathbb{R}$ that evaluates the interpretability of a strategy (σ) in a game environment (\mathcal{G}) by assigning it a numerical score.

◀ Return

Divergence

Move Evaluation Function

Given chess engine E with position evaluation function $v_E(s)$, we can obtain a move evaluation function $q_E(s, a)$ as —

$$q_E(s, a) = \sum_{s', r} \mathcal{P}(s', r | s, a) [r + v_E(s')] \quad (4)$$

$$= v_E(s'), s' \text{ is non-terminal} \quad (5)$$

Equation 5 follows from 4 since rewards in chess are 0 for non-terminal states, $\gamma = 1$, and chess rules are deterministic.

Divergence

Difference Function

Given two moves a_1, a_2 made in a position s , we can calculate their difference $d_E(s, a_1, a_2)$ as —

$$d_E(s, a_1, a_2) \doteq | q_E(s, a_1) - q_E(s, a_2) | \quad (6)$$

◀ Return

Metrics

Definition

Coverage is the fraction of positions in a given set P to which the chess strategy model is *applicable*.

$$P_A \doteq P \cap A_\sigma \quad (7)$$

$$\text{Coverage}(\sigma, P) \doteq \frac{|P_A|}{|P|} \quad (8)$$

Divergence

Definition (Divergence)

Divergence of a tactic from a set of examples P is the average difference in *evaluation* between the moves suggested by the tactic and the ground truth move.

$$\text{Divergence}_E(\sigma, P) \doteq \frac{1}{|P_A|} \sum_{(s, a_1) \in P_A} \sum_{a_2 \in \mathcal{A}(s)} \sigma(a_2 | s) d_E(s, a_1, a_2) \quad (9)$$

◀ Return

PAL

- Patterns and Learning (Morales, 1992)
- ILP system to learn chess *patterns*
- Predicate vocabulary
- *rllg* algorithm + heuristics to learn patterns
- Automatic *example generator* to learn target concepts

◀ Return

Precision/Recall-based Constraints

Definition (Precision constraint)

A precision constraint prunes the specializations of a hypothesis if its precision on a set of examples is less than some pre-defined lower limit.

Definition (Recall constraint)

A recall constraint prunes specializations of a hypothesis if its recall on a set of examples is less than some pre-defined lower limit.

◀ Return

Precision/Recall-based Constraints

Theorem

Given hypotheses $H_1, H_2 \in \mathbb{H}$ with $H_1 \preceq H_2$ and having recall values of r_1 and r_2 on a training set respectively, then $r_1 \leq r_2$.

◀ Return

Precision/Recall-based Constraints

Proof.

$$\text{Recall} = \frac{tp}{tp + fn} \quad (10)$$

$$H_1 \preceq H_2 \quad (11)$$

$$r_1 = \frac{tp_1}{tp_1 + fn_1} \quad (12)$$

$$r_2 = \frac{tp_2}{tp_2 + fn_2} \quad (13)$$

By the definition of theory subsumption, we have

$$tp_1 \leq tp_2, \text{ and} \quad (14)$$

$$fp_1 \leq fp_2 \quad (15)$$

Since the number of positive examples is constant for the same population,

$$\therefore r_1 \leq r_2 \quad (16)$$

[◀ Return](#)

Predicate Vocabulary

- Allows more *situational rule* expression – en passant, promotion
- Allows *more efficient* unification due to introduction of more types
- Adds higher-order state features like *can_threat* and *xrays*

◀ Return

Interpretability Measure

Interpretability Measure

- *No explicit interpretability measure!* Only qualitative arguments

Interpretability Measure

- *No explicit interpretability measure!* Only qualitative arguments
- Human players *think* and *train* using chess tactics (Gobet & Jansen, 2006;
Szabo, 1984)

Interpretability Measure

- *No explicit interpretability measure!* Only qualitative arguments
- Human players *think* and *train* using chess tactics (Gobet & Jansen, 2006; Szabo, 1984)
- FO-logic used extensively to model chess patterns (Berliner, 1975; Bramer, 1977; Bratko, 1982; Huberman, 1968; Morales, 1992; Pitrat, 1977; Wilkins, 1979)

Interpretability Measure

- *No explicit interpretability measure!* Only qualitative arguments
- Human players *think* and *train* using chess tactics (Gobet & Jansen, 2006; Szabo, 1984)
- FO-logic used extensively to model chess patterns (Berliner, 1975; Bramer, 1977; Bratko, 1982; Huberman, 1968; Morales, 1992; Pitrat, 1977; Wilkins, 1979)
- Logic rules are *acknowledged to be interpretable* (Zhang et al., 2021)

Answer Set Programming

- *Declarative programming* paradigm based on *stable models* (Gelfond & Lifschitz, 1988)
- ASP language (Gebser, Harrison, Kaminski, Lifschitz, & Schaub, 2015) allows using rules to —
 - *model* a design space
 - *restrict* it using integrity constraints
 - *generate* instances in the newly restricted space

◀ Return

Example

```
1 #const width=10.
2
3 param("width",width).
4
5 dim(1..width).
6
7 tile((X,Y)) :- dim(X), dim(Y).
8
9 adj((X1,Y1),(X2,Y2)) :- tile((X1,Y1)), tile((X2,Y2)), \
0   #abs(X1-X2)+#abs(Y1-Y2) == 1.
1
2 start((1,1)). finish((width,width)).
3
4 % tiles have at most one named sprite
5 0 { sprite(T,wall;gem;altar) } 1 :- tile(T).
6
7 % there is exactly one altar and one gem in the whole level
8 :- not 1 { sprite(T,altar) } 1. :- not 1 { sprite(T,gem) } 1.
```

An ASP program which can generate maze-like levels with integrity constraints that specify the number of game objects.

Transformer

- Neural-network based model
- Great success at learning sequences
- Uses an *attention* mechanism to focus on relevant parts of a sequence

$$z_i = \sum_{j=1}^n \text{softmax}(\{\langle q_i, k_{j'} \rangle\}_{j'=1}^n)_j \cdot v_j \quad (17)$$

Rise of Esports



Colin Young-Wolff/Riot Games

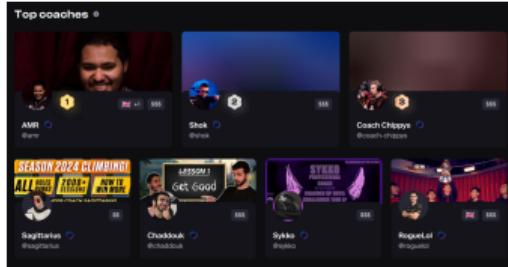
T1 huddles with the trophy onstage after their victory at the League of Legends World Championship 2023 Finals at the Gocheok Sky Dome on November 19, 2023 in Seoul, South Korea.

Coaching in Esports



Robert Paul

Professional *Tekken 7* player Arslan Ash receives on-stage coaching during the *Tekken World Tour 2023 Finals*



Metafy

Top *League of Legends* coaches available for hire on the Metafy platform



RyderDie

A beginner guide for *Counter Strike 2* on YouTube

Difficulties with Esports Coaching

Esports coaching is –

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods
 - learning resources are *scattered* and *unstructured*

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods
 - learning resources are *scattered* and *unstructured*
- *confined* to popular games

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods
 - learning resources are *scattered* and *unstructured*
- *confined* to popular games
- *difficult* due to a shifting “*meta*”

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods
 - learning resources are *scattered* and *unstructured*
- *confined* to popular games
- *difficult* due to a shifting “*meta*”
- not *scalable*

Difficulties with Esports Coaching

Esports coaching is –

- not *standardized*: no “textbook” to follow
 - coaches follow *ad-hoc* methods
 - learning resources are *scattered* and *unstructured*
- *confined* to popular games
- *difficult* due to a shifting “*meta*”
- not *scalable*

Solution

A *principled* and *automated* approach to esports coaching could democratize access to it for all interested players.

Summary

Research Thrust	RQ	Sub-RQ	Publication
ISS Framework	RQ1	–	EAAI '22 <small>(Krishnan & Martens, 2022b)</small>
		RQ2(a)	
ISS using Logic	RQ2	RQ2(b)	SG+EA Workshop @ AIIDE '22 <small>(Krishnan & Martens, 2022a)</small>
		RQ2(c)	
ISS using Programs	RQ3	RQ3(a)	ACS '24 <small>(Krishnan et al., 2024)</small>
		RQ3(b)	

Summary

Research Thrust	RQ	Sub-RQ	Publication
<i>ISS Framework</i>	RQ1	–	EAAI '22 <small>(Krishnan & Martens, 2022b)</small>
		RQ2(a)	
ISS using Logic	RQ2	RQ2(b)	SG+EA Workshop @ AIIDE '22 <small>(Krishnan & Martens, 2022a)</small>
		RQ2(c)	
ISS using Programs	RQ3	RQ3(a)	ACS '24 <small>(Krishnan et al., 2024)</small>
		RQ3(b)	

Summary

Research Thrust	RQ	Sub-RQ	Publication
ISS Framework	RQ1	–	EAAI '22 <small>(Krishnan & Martens, 2022b)</small>
		RQ2(a)	
<i>ISS using Logic</i>	RQ2	RQ2(b)	SG+EA Workshop @ AIIDE '22 <small>(Krishnan & Martens, 2022a)</small>
		RQ2(c)	
ISS using Programs	RQ3	RQ3(a)	ACS '24 <small>(Krishnan et al., 2024)</small>
		RQ3(b)	

Summary

Research Thrust	RQ	Sub-RQ	Publication
ISS Framework	RQ1	–	EAAI '22 <small>(Krishnan & Martens, 2022b)</small>
		RQ2(a)	
ISS using Logic	RQ2	RQ2(b)	SG+EA Workshop @ AIIDE '22 <small>(Krishnan & Martens, 2022a)</small>
		RQ2(c)	
<i>ISS using Programs</i>	RQ3	RQ3(a)	ACS '24 <small>(Krishnan et al., 2024)</small>
		RQ3(b)	

Why Chess?

Why Chess?

- *Popular* game with a *long* competitive history

Why Chess?

- *Popular* game with a *long* competitive history
- Has a large number of *player-discovered strategies*

Why Chess?

- *Popular* game with a *long* competitive history
- Has a large number of *player-discovered strategies*
- Extensive use as a *testbed for AI*
 - Convenient software implementations
 - Work in learning chess patterns
 - *Engines* for evaluation

Dataset

- Randomly generated program dataset (Trivedi et al., 2021)
- Convert programs into a trajectory in the program synthesis MDP

Transitions	870,782
Programs	50,000
Program Length	19.30 ± 5.96 (min = 7, max = 50)
Episode Length	17.42 ± 4.33 (min = 5, max = 30)

Table: Statistics for the LEAPS program dataset.

Limitations and Future Work

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required
~ 800 LoC of Prolog!

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required
~ 800 LoC of Prolog!
 - *general* policy BK

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required ~ 800 LoC of Prolog!
 - *general* policy BK
- Performance not competitive with state-of-the-art

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required ~ 800 LoC of Prolog!
 - *general* policy BK
- Performance not competitive with state-of-the-art
 - experiment with *task-specific* BK

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required ~ 800 LoC of Prolog!
 - *general* policy BK
- Performance not competitive with state-of-the-art
 - experiment with *task-specific* BK
 - use *neural* ILP systems (Evans & Grefenstette, 2018)

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required ~ 800 LoC of Prolog!
 - *general* policy BK
- Performance not competitive with state-of-the-art
 - experiment with *task-specific* BK
 - use *neural* ILP systems (Evans & Grefenstette, 2018)
- Interpretability for FOL-based strategies is poorly understood

Limitations and Future Work

- High cost of *knowledge engineering* – improved BK required ~ 800 LoC of Prolog!
 - *general* policy BK
- Performance not competitive with state-of-the-art
 - experiment with *task-specific* BK
 - use *neural* ILP systems (Evans & Grefenstette, 2018)
- Interpretability for FOL-based strategies is poorly understood
 - study factors affecting interpretability (Kliegr et al., 2021)

Limitations and Future Work

Limitations and Future Work

- Different DSLs required for different domains

Limitations and Future Work

- Different DSLs required for different domains
 - *general* policy DSL (Qiu & Zhu, 2022; Verma et al., 2018)

Limitations and Future Work

- Different DSLs required for different domains
 - *general* policy DSL (Qiu & Zhu, 2022; Verma et al., 2018)
- Interpretability for programmatic strategies is poorly understood

Limitations and Future Work

- Different DSLs required for different domains
 - *general* policy DSL (Qiu & Zhu, 2022; Verma et al., 2018)
- Interpretability for programmatic strategies is poorly understood
 - study factors affecting interpretability of programmatic strategies

Expected Outcomes

- Benefit esports industry → *better coaching* for players
- Benefit explainable AI research → generate *policy explanations*

Background and Related Work

Automated Coaching for Esports

Background and Related Work

Automated Coaching for *Esports*

- Esports Analytics
- Automated Game Analysis

Esports Analytics

Esports Analytics

- *Data-driven* analysis of esports games to provide *support* to players

Esports Analytics

- *Data-driven* analysis of esports games to provide *support* to players
- win prediction (Clark et al., 2020; Hodge et al., 2021; Maymin, 2021; Novak et al., 2020; Schubert et al., 2016)

Esports Analytics

- *Data-driven* analysis of esports games to provide *support* to players
- win prediction (Clark et al., 2020; Hodge et al., 2021; Maymin, 2021; Novak et al., 2020; Schubert et al., 2016)
- patch-agnostic methods (Chitayat et al., 2023)

Esports Analytics

- *Data-driven* analysis of esports games to provide *support* to players
- win prediction (Clark et al., 2020; Hodge et al., 2021; Maymin, 2021; Novak et al., 2020; Schubert et al., 2016)
- patch-agnostic methods (Chitayat et al., 2023)
- eye-tracking (Nagel, 2017; Stepanov et al., 2019)

Esports Analytics

- *Data-driven* analysis of esports games to provide *support* to players
- win prediction (Clark et al., 2020; Hodge et al., 2021; Maymin, 2021; Novak et al., 2020; Schubert et al., 2016)
- patch-agnostic methods (Chitayat et al., 2023)
- eye-tracking (Nagel, 2017; Stepanov et al., 2019)
- data collection (Frommel et al., 2019; Xenopoulos et al., 2022)

Automated Game Analysis

Automated Game Analysis

- automated *testing* to discover bugs (Bergdahl et al., 2020; Pfau et al., 2017; Zheng et al., 2019)

Automated Game Analysis

- automated *testing* to discover bugs (Bergdahl et al., 2020; Pfau et al., 2017; Zheng et al., 2019)
- automated *game balancing* to ensure fairness (Beau & Bakkes, 2016; DeLaurentis et al., 2021)

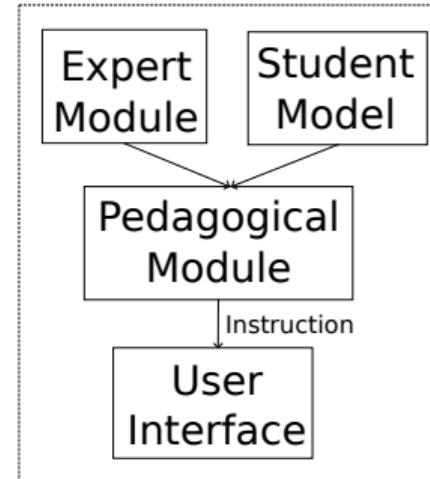
Background and Related Work

Automated *Coaching* for Esports

- Intelligent Tutoring System
- Player Modeling

Intelligent Tutoring System (ITS)

- *Automated, personalized* education system
- No ITS for esports (Mousavinasab et al., 2021)



ITS architecture

Player Modeling

- To provide *personalized* experiences in games
 - dynamic difficulty adjustment (Xue, Wu, Kolen, Aghdaie, & Zaman, 2017)
 - generated levels (Jennings-Teats, Smith, & Wardrip-Fruin, 2010)
- Strategy = player model (Smith, Lewis, Hullet, Smith, & Sullivan, 2011)

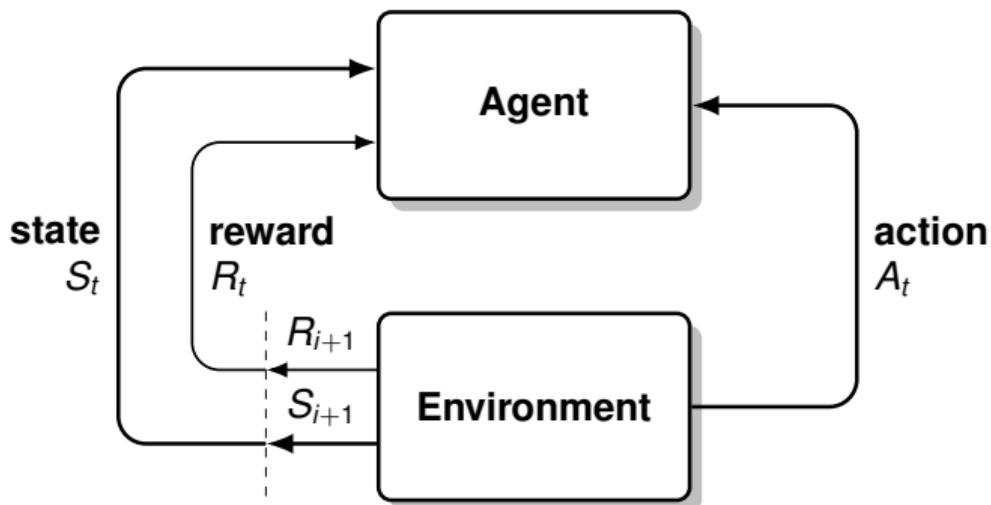
Background and Related Work

Automated Coaching for Esports

- Reinforcement Learning (RL)
- Explainable RL
- Strategy Synthesis

Reinforcement Learning (RL)

- ML paradigm where an agent *learns to act* in an environment
- Game strategies = RL *policies*



The agent-environment interaction interface in a MDP. Reproduced from Figure 3.1 of Sutton and Barto (2018).

Explainable Reinforcement Learning (XRL)

Explainable Reinforcement Learning (XRL)

- Modern RL agents are *difficult* to understand compared to earlier approaches

Explainable Reinforcement Learning (XRL)

- Modern RL agents are *difficult* to understand compared to earlier approaches
- Agents are being used in *high-stakes* applications

Explainable Reinforcement Learning (XRL)

- Modern RL agents are *difficult* to understand compared to earlier approaches
- Agents are being used in *high-stakes* applications
- Necessary to *interpret* an agent and *explain* their decisions

Explainable Reinforcement Learning (XRL)

- Modern RL agents are *difficult* to understand compared to earlier approaches
- Agents are being used in *high-stakes* applications
- Necessary to *interpret* an agent and *explain* their decisions
- XRL: field of study to devise techniques to interpret RL agents

Explainable Reinforcement Learning (XRL)

- Modern RL agents are *difficult* to understand compared to earlier approaches
- Agents are being used in *high-stakes* applications
- Necessary to *interpret* an agent and *explain* their decisions
- XRL: field of study to devise techniques to interpret RL agents
- ISS = XRL for game environments

Strategy Synthesis

Strategy Synthesis

- Early work uses rule-based models (Butler et al., 2017; Canaan et al., 2018; de Freitas et al., 2018;

de Mesentier Silva et al., 2016; Spronck et al., 2004)

Strategy Synthesis

- Early work uses rule-based models (Butler et al., 2017; Canaan et al., 2018; de Freitas et al., 2018; de Mesentier Silva et al., 2016; Spronck et al., 2004)
- Subsequent work uses programmatic representations (Mariño et al., 2021; Mariño & Toledo, 2022; Medeiros et al., 2022)

Strategy Synthesis

- Early work uses rule-based models (Butler et al., 2017; Canaan et al., 2018; de Freitas et al., 2018; de Mesentier Silva et al., 2016; Spronck et al., 2004)
- Subsequent work uses programmatic representations (Mariño et al., 2021; Mariño & Toledo, 2022; Medeiros et al., 2022)
- No unifying framework to synthesize goals and approaches